



**University of
Zurich**^{UZH}

**Understanding Data Mixture Effects in Financial Language Model
Pretraining**
A Study of Domain-Specific and High-Quality General Corpora

MASTER'S THESIS

SUBMITTED IN PARTIAL FULFILLMENT OF THE REQUIREMENTS FOR THE DEGREE OF MASTER OF
ARTS IN ECONOMICS AND BUSINESS ADMINISTRATION

STUDENT
GUANLAN LIU
19-768-837
GUANLAN.LIU@UZH.CH

SUPERVISOR
PROF. DR. MARKUS LEIPPOLD
PROFESSOR OF FINANCIAL ENGINEERING
DEPARTMENT OF FINANCE
UNIVERSITY OF ZURICH

MIN YANG
MIN.YANG2@UZH.CH

DATE OF SUBMISSION: [DATE]

Abstract

We present a compute-normalized study of pretraining data composition for financial language models. We adapt the decoder-only Qwen3 Base architecture to a fixed 100M-token budget with heterogeneous financial texts and general texts with a unified eight-dataset evaluation. We further develop systematic comparisons of individual datasets versus mixtures to evaluate optimal pretraining strategies. In particular, we find that **medium individual datasets (3.6–8.5M tokens) consistently outperform mixtures on both performance and consistency**. FiQA (6.80 ppl, 19% spread), FinGPT (7.03 ppl, 37% spread), and Alpaca (8.73 ppl, 11.5% spread) achieve $2.5\text{--}3.2\times$ better perplexity AND $1.5\text{--}4.8\times$ better cross-dataset consistency than our seven-source financial mixture (21.55 ppl, 55% spread). This finding challenges conventional wisdom that data diversity improves robustness. This occurs through a three-way interaction: medium datasets achieve optimal epoch counts (12–28 epochs) with format consistency, while large datasets undertrain (≤ 1 epoch) and large mixtures add format conflicts that small models (0.6B–4B) cannot reconcile. Small datasets (≥ 1 M tokens) overtrain (143–352 epochs), leading to memorization. WikiText shows competitive performance at small scales (0.6B: 9.68 ppl) but reverse scaling at larger sizes due to training instability.

Our contributions in this work are three-folded:

- a. We systematically compare individual datasets versus mixtures via token-matched training and unified eight-dataset evaluation, revealing that medium individual datasets (3.6–8.5M tokens) consistently outperform mixtures on both performance and consistency metrics.
- b. To understand why mixtures fail, we analyze format inconsistency, vocabulary dilution, and multi-task interference effects. We find that focused optimization on single datasets beats diverse mixing—format consistency and concentrated vocabulary exposure outweigh anticipated diversity benefits.
- c. We establish that data quality and focus matter more than scale: medium datasets (FiQA 3.6M, FinGPT 4.1M, Alpaca 8.5M, SEC 8.1M) substantially outperform large datasets (News 194M, WikiText 124M). This non-monotonic size-performance relationship reflects optimal epoch counts (12–28) combined with format consistency, outperforming both undertrained large datasets (≤ 1 epoch) and overtrained small datasets (143–352 epochs).

Contents

1	Introduction	1
1.1	Motivation	1
1.2	Research Questions	1
1.3	Related Work	2
1.4	Contributions	3
2	Background and Related Work	4
2.1	Financial NLP	4
2.1.1	Tasks in Financial NLP	4
2.1.2	Existing Financial Language Models	4
2.1.3	Domain-Specific Challenges	4
2.2	Language Model Pretraining	5
2.2.1	Pretraining Objectives and Architecture	5
2.2.2	Scaling Laws and Model Size Effects	5
2.2.3	Computational and Memory Considerations	5
2.3	Data Mixture Strategies	5
2.3.1	Curriculum Learning and Sequential Mixing	5
2.3.2	Simultaneous Mixture Approaches	6
2.3.3	Domain Proportions and Sampling Strategies	6
2.4	Domain Adaptation and Transfer Learning	6
2.4.1	Cross-Domain Transfer in Language Models	6
2.4.2	Catastrophic Forgetting and Stability	7
2.4.3	Distribution Shift and Domain Mismatch	7
2.4.4	Related Empirical Studies	7
3	Methodology	8
3.1	Experimental Design Overview	8
3.2	Model Architecture	9
3.3	Datasets	9
3.3.1	Financial Datasets	9

3.3.2	WikiText	10
3.3.3	Mixture Strategies	10
3.4	Training Setup and Hyperparameter Tuning	12
3.4.1	Initial Configuration	12
3.4.2	Pragmatic Learning Rate Adjustments	12
3.4.3	Other Hyperparameters	12
3.4.4	Computational Budget	12
3.5	Evaluation Protocol	13
3.5.1	Multi-Dataset Evaluation	13
3.5.2	Metrics	13
4	Results	14
4.1	Overview of Experimental Results	14
4.2	Individual Datasets vs Mixtures	15
4.2.1	Mixed Financial Datasets: Inferior on All Metrics	15
4.2.2	Mixed Wiki+Financial	16
4.2.3	Pure WikiText Baseline	18
4.2.4	Key Takeaways	19
4.3	Individual Dataset Analysis: Component Effects	20
4.3.1	Large Datasets	20
4.3.2	Medium Datasets	20
4.3.3	Small Datasets	22
4.3.4	Dataset Size vs Generalization	24
4.4	Training Dynamics and Scaling Behavior	26
4.4.1	Normal Scaling Pattern	26
4.4.2	Reverse Scaling Phenomenon	28
4.4.3	Learning Rate Sensitivity by Model Size	29
4.4.4	Fixing Reverse Scaling	30
4.4.5	Model Stability Analysis	30
4.5	Domain Transfer and Generalization Patterns	31
4.5.1	Cross-Dataset Evaluation	31
4.5.2	Document Format and Task Type Effects	32
4.5.3	Variance Comparison	34
4.5.4	Domain-Specific vs General Knowledge Transfer	35
4.6	Summary and Key Results	37
5	Discussion	43
5.1	Key Empirical Findings	43
5.2	Interpretation of Data Interaction Effects	44

5.2.1	Why WikiText Underperforms on Financial Tasks	44
5.2.2	Why Mixtures Fail: The Diversity Paradox	45
5.2.3	The Optimal Training Regime: Why Medium Datasets Dominate	46
5.2.4	Domain Interference Patterns	46
5.2.5	Scale-Dependent Training Notes	47
5.3	Practical Guidelines for Financial LM Pretraining	47
5.3.1	Data Mixture Strategies by Use Case	47
5.3.2	Model Size Selection	48
5.3.3	Learning Rate Notes	49
5.3.4	Token Budget Allocation	49
5.4	Limitations and Threats to Validity	49
6	Conclusion	51
6.1	Summary of Contributions	51
6.1.1	Individual Datasets Outperform Mixtures	51
6.1.2	Learning Rate Notes	52
6.1.3	Dataset Size Effects and Generalization	52
6.1.4	Domain Transfer and Format Effects	52
6.1.5	Model Size Selection for Resource-Constrained Settings	53
6.1.6	Open-Source Reproducible Pipeline	53
6.2	Implications for Practice and Research	54
6.2.1	Actionable Deployment Guidelines	54
6.2.2	For Researchers: Open Questions and Methodological Lessons	54
6.2.3	For Industry: Privacy-Preserving Financial AI	55
6.3	Future Research Directions	55
6.3.1	Scaling to Larger Models and Architectures	55
6.3.2	Advanced Mixture Optimization	56
6.3.3	Comprehensive Downstream Evaluation	56
6.3.4	Multi-Stage Pretraining Strategies	56
6.3.5	Open Questions	57
6.4	Closing Remarks	57

List of Figures

3.1	50cap Mixture Strategy Visualization	11
4.1	Mixed Financial Dataset: Scaling Behavior	16
4.2	Mixed Wiki+Financial Dataset: Scaling Behavior	17
4.3	WikiText Dataset: Reverse Scaling	18
4.4	Comparison of Mixture Strategies	20
4.5	Financial News Dataset: Scaling Behavior	21
4.6	SEC Reports Dataset: Scaling Behavior	21
4.7	FinGPT Sentiment Dataset: Scaling Behavior	23
4.8	Finance Alpaca Dataset: Scaling Behavior	23
4.9	FiQA Dataset: Scaling Behavior	24
4.10	Financial QA 10K Dataset: Reverse Scaling	26
4.11	Twitter Financial Sentiment Dataset: Reverse Scaling	26
4.12	Dataset Size vs. Generalization Across Model Scales	28
4.13	Cross-Dataset Transfer Heatmap	41
4.14	Cross-Dataset Variance Comparison	42

List of Tables

3.1	Experimental Settings Summary	8
3.2	Qwen3 Model Specifications	9
3.3	Financial Dataset Characteristics	10
3.4	WikiText Dataset Characteristics	10
4.1	Overview of Pretraining Experiments	14
4.2	Mixed Financial: Evaluation Results	16
4.3	Mixed Wiki+Financial: Evaluation Results	17
4.4	WikiText: Learning Rate Comparison	19
4.5	Financial News: Evaluation Results	22
4.6	SEC Reports: Evaluation Results	22
4.7	FinGPT Sentiment: Evaluation Results	24
4.8	Finance Alpaca: Evaluation Results	25
4.9	FiQA: Evaluation Results	25
4.10	Financial QA 10K: Learning Rate Comparison	27
4.11	Twitter Financial: Learning Rate Comparison	27
4.12	Financial News Evaluation: Cross-Dataset Performance	32
4.13	SEC Reports Evaluation: Cross-Dataset Performance	33
4.14	Alpaca Evaluation: Cross-Dataset Performance	34
4.15	FinGPT Evaluation: Cross-Dataset Performance	35
4.16	FiQA Evaluation: Cross-Dataset Performance	36
4.17	Twitter Financial Evaluation: Cross-Dataset Performance	37
4.18	Financial QA Evaluation: Cross-Dataset Performance	38
4.19	WikiText Evaluation: Cross-Dataset Performance	39
4.20	Best Configurations by Application	39

Chapter 1

Introduction

1.1 Motivation

Large language models (LLMs) have rapidly changed how we do natural language processing (Vaswani et al. 2017; Radford et al. 2019; Brown et al. 2020; Touvron et al. 2023). Yet using them in finance still brings practical hurdles. Financial institutions and individuals handle highly sensitive data—transactions, portfolios, trading strategies—that cannot be sent to external APIs for privacy and competitive reasons (e.g., GDPR) (*Regulation (EU) 2016/679 of the European Parliament and of the Council of 27 April 2016 on the protection of natural persons with regard to the processing of personal data and on the free movement of such data (General Data Protection Regulation)* 2016). We therefore need lightweight, locally runnable financial language models that maintain reasonable performance while protecting data.

In practice, domain adaptation tends to follow two paths: train very large models from scratch or fine-tune general models on domain data. Most teams cannot afford the first; the second often misses domain nuances (Gururangan et al. 2020). And there is common practice showing that high quality general corpora (e.g., Wikipedia, The Pile) always help specialized applications. However, researchers also show that exceptions do appear (Gao et al. 2021; Raffel et al. 2020; Longpre et al. 2023).

This thesis studies how different data sources, both in-domain financial data and out-of-domain high-quality corpora, interact during pretraining. We focus on models in the 0.6B to 4B parameter range, which are realistic for laptops and some mobile devices while keeping acceptable performance (A. Yang et al. 2024; Xia et al. 2023; Team et al. 2024; Javaheripi et al. 2023). Through systematic experiments across 10 pretraining configurations and three model sizes, we present evidence about data mixture strategies for specialized domains.

1.2 Research Questions

This thesis investigates the following core research questions.

RQ1: Individual datasets versus mixtures Do data mixtures improve performance and consistency compared to individual dataset training? Contrary to conventional wisdom, our results show that **medium individual datasets (3.6–8.5M tokens) strictly dominate mixtures on both metrics**. FiQA (6.80 ppl, 19% spread), FinGPT (7.03 ppl, 37% spread), and Alpaca (8.73 ppl,

11.5% spread) achieve $2.5\text{--}3.2\times$ better perplexity AND $1.5\text{--}4.8\times$ better cross-dataset consistency than Mixed Financial (21.55 ppl, 55% spread). This finding (Figure 4.4 and Tables 4.2 and 4.7 to 4.9) challenges the mixture hypothesis—data diversity degrades both performance and robustness at fixed token budgets (100M).

RQ2: Model size and training dynamics How do optimal training configurations vary across model sizes (0.6B/1.7B/4B parameters)? What is the relationship between size and learning rate sensitivity? In our setup, we trained all main runs with LR=2e-5; for a few abnormal cases, we reduced LR pragmatically (e.g., to 1×10^{-5} or 5×10^{-6}) and saw improved stability.

RQ3: Dataset size and quality effects What is the relationship between dataset size and performance? Surprisingly, we find a **non-monotonic relationship**: medium datasets (3.6–8.5M tokens) achieve the best results, outperforming both small ($\approx 1\text{M}$) and large ($\gtrsim 100\text{M}$) datasets. Medium datasets FiQA (3.6M, 6.80 ppl), FinGPT (4.1M, 7.03 ppl), Alpaca (8.5M, 8.73 ppl), and SEC (8.1M, 17.80 ppl) substantially beat the large dataset News (194M, 32.82 ppl). This suggests data quality, focus, and format consistency matter more than scale. Medium datasets achieve optimal training (12–28 epochs per dataset) with format consistency, enabling focused learning. Large datasets undertrain (≈ 1 epoch), while large mixtures combine undertraining with format inconsistency that small models (0.6B–4B) struggle to resolve. Small datasets overtrain (143–352 epochs), causing memorization. Only very small datasets ($\approx 1\text{M}$ tokens) exhibit severe overtraining (Figures 4.10 and 4.11), while medium datasets achieve optimal performance without mixing.

RQ4: Domain transfer patterns How well do models pretrained on financial data transfer across task types (sentiment, question answering, document understanding), and how much does document format matter? Cross dataset comparison tables (Tables 4.12 to 4.17) suggest that format consistency (long form, instruction, short form) drives transfer more than domain vocabulary, with boldface patterns clustering along format based diagonals.

These questions are addressed through a detailed experimental framework with more than 30 trained models and 230 evaluation results across eight held-out test sets, providing systematic evidence on data mixture effects in specialized-domain pretraining.

1.3 Related Work

Financial NLP. The domain covers sentiment analysis, question answering, numerical reasoning, and information extraction from regulatory documents (Araci 2019; Z. Chen et al. 2021). Challenges include specialized vocabulary (“alpha,” “EBITDA”), domain reasoning patterns, and privacy constraints that push toward local deployment (S. Wu et al. 2023). Existing models range from FinBERT variants (Araci 2019; Y. Yang et al. 2020) to BloombergGPT (50B, mixed 51% financial and 49% general) (S. Wu et al. 2023). Most focus on single large models, while few study mixture effects across sizes.

Language model pretraining. Modern LMs use causal language modeling (next-token prediction) on transformer architectures (Vaswani et al. 2017; Radford et al. 2019; Brown et al. 2020). Scaling laws (J. Kaplan et al. 2020; Hoffmann et al. 2022) show that model size, data size, and compute follow power-law relationships. Bigger models are more sample-efficient. But hyperparameter sensitivity at intermediate scales (0.6B–4B) is less studied. Chapter 2 covers training dynamics and memory constraints.

Data mixture strategies. Pretraining can be sequential (curriculum) or simultaneous. In our study, we use a simple “50cap” mixture and refer to Chapter 2 for rationale.

Domain adaptation and transfer. Transfer learning assumes general pretraining helps specialized tasks (Devlin et al. 2019; Pan and Q. Yang 2010). But Gururangan et al. (2020) showed domain-adaptive pretraining (continued training on domain text) improves performance. Key challenges include catastrophic forgetting (Kirkpatrick et al. 2017) and distribution shift (Quiñonero-Candela et al. 2008).

1.4 Contributions

This thesis makes six primary contributions that challenge conventional assumptions about data mixture strategies for language model pretraining.

First, we overturn the mixture hypothesis: **medium individual datasets (3.6–8.5M tokens) consistently outperform mixtures on both performance and consistency.** FiQA (6.80 ppl, 19% spread), FinGPT (7.03 ppl, 37% spread), and Alpaca (8.73 ppl, 11.5% spread) achieve 2.5–3.2 \times better perplexity AND 1.5–4.8 \times better cross-dataset consistency than Mixed Financial (21.55 ppl, 55% spread). This finding contradicts widespread assumptions that data diversity improves robustness. We explain why mixtures fail through analysis of format inconsistency, vocabulary dilution, and multi-task interference, supported by 11 scaling figures and 18 detailed evaluation tables.

Second, we establish a **non-monotonic relationship between dataset size and performance:** medium datasets (3.6–8.5M) substantially outperform large datasets ($> 100M$). FiQA (3.6M), FinGPT (4.1M), and Alpaca (8.5M) beat News (194M) by 2–5 \times on average perplexity, suggesting data quality, focus, and format consistency matter more than scale. SEC (8.1M) also performs well as a medium dataset. Only very small datasets ($< 1M$) exhibit severe overtraining and require careful tuning or exclusion.

Third, we document pragmatic learning-rate adjustments. All main runs used LR=2e-5; in three cases with training instabilities (WikiText, Financial QA, Twitter), reducing LR (to 1×10^{-5} or 5×10^{-6}) stabilized optimization. These are empirical fixes, not theoretical scaling laws.

Fourth, we analyze domain transfer patterns, finding format consistency drives transfer more than vocabulary overlap. Long-form documents transfer well (News \leftrightarrow SEC), instruction tasks cluster (FinGPT/Alpaca/FiQA), but cross-format transfer fails despite shared domain.

Fifth, we demonstrate WikiText shows scale-dependent behavior: competitive at 0.6B (9.68 ppl) but reverse scaling at 4B (31.54 ppl) due to training instability, not domain mismatch. Adding WikiText to financial mixtures degrades performance (26.69 ppl vs 21.55 for pure financial).

Lastly, we show lightweight financial models (0.6B–4B) deliver practical performance for edge deployment when trained on focused medium datasets, enabling privacy-preserving financial NLP without external API dependencies.

Chapter 2

Background and Related Work

We review the literature in this chapter. We start with financial NLP, then pretraining basics, prior mixture strategies, and finally domain adaptation and transfer learning.

2.1 Financial NLP

2.1.1 Tasks in Financial NLP

Financial natural language processing covers many tasks: sentiment on news and social media, question answering on regulatory documents, numerical reasoning in reports, and information extraction from SEC filings (Araci 2019; Z. Chen et al. 2021). The domain has specific challenges compared to general NLP: specialized vocabulary (e.g., “alpha”, “beta”, “EBITDA”), domain reasoning patterns (e.g., causal chains in market analysis), numerical grounding (financial statements), and temporal dynamics (market events, earnings releases) (S. Wu et al. 2023; Araci 2019).

2.1.2 Existing Financial Language Models

Several finance focused language models appeared in recent years. **BloombergGPT** (S. Wu et al. 2023), a 50 billion parameter model, was pretrained on a mixture of 51% financial and 49% general data, showing strong performance on financial benchmarks while keeping general capabilities. **FinBERT** variants (Araci 2019; Y. Yang et al. 2020) adapted BERT to financial text via continued pretraining, improving sentiment analysis on financial news. More recently, **FinGPT** (H. Yang et al. 2023) explored open source instruction tuning for financial tasks. Together, these work show both scale first and adaptation first approaches.

2.1.3 Domain-Specific Challenges

Financial NLP faces three practical challenges. First, privacy concerns: financial institutions cannot upload sensitive data (portfolios, trading strategies, client information) to external APIs, so locally deployable models are needed (S. Wu et al. 2023). Second, data scarcity: compared to general web text, curated financial corpora are smaller, so data-efficient training is important. Third, rapid vocabulary change: financial language shifts with market trends (e.g., “DeFi”, “ESG”), so models

must adapt to new terms. These constraints motivate our focus on 0.6B to 4B models, as we believe larger models should have better generalization and adaptation capabilities.

2.2 Language Model Pretraining

2.2.1 Pretraining Objectives and Architecture

Modern language models mostly use **causal language modeling**: predict the next token from the context (Radford et al. 2019; Brown et al. 2020). We follow this default. Architecturally, we use the usual decoder-only transformer (GPT, LLaMA, Qwen): self-attention for long context and feed-forward blocks for the non-linear part (Vaswani et al. 2017; Touvron et al. 2023).

2.2.2 Scaling Laws and Model Size Effects

The work of J. Kaplan et al. (2020) linked model size, dataset size, compute, and final performance by power laws. The core point, that larger models can be more sample efficient, pushed the field toward billion parameter scales. Later work added nuance: Hoffmann et al. (2022) showed undertraining is common (Chinchilla); Tay et al. (2022) emphasized objectives and data quality.

2.2.3 Computational and Memory Considerations

Training large language models requires substantial compute. A 1 billion parameter model with 32 bit precision uses roughly 4GB of memory for parameters alone, with optimizer states (e.g., Adam’s momentum terms) doubling or tripling this requirement (Rajbhandari et al. 2020; Kingma and Ba 2014).

For models in the 0.6B to 4B range targeted here, memory efficient techniques like mixed precision (bfloating16), gradient accumulation, activation checkpointing, and parameter efficient fine tuning such as LoRA allow training on enterprise class GPUs (e.g., NVIDIA RTX A6000 48GB, A100 40GB, H100 80GB) (Narayanan et al. 2021; Hu et al. 2021). In practice, these techniques are very important for us as we rely on renting GPUs from lambda¹ and we are striving to save compute.

2.3 Data Mixture Strategies

2.3.1 Curriculum Learning and Sequential Mixing

Curriculum learning in language model pretraining involves carefully sequencing training data from easier to harder examples, or from general to specialized domains (Bengio et al. 2009). Zhang et al. (2022) applied curriculum strategies in pretraining OPT models, progressively increasing data difficulty. In the financial domain, a natural curriculum might proceed from general Wikipedia text to financial news to technical SEC filings. However, empirical evidence for curriculum’s effectiveness in large-scale pretraining remains mixed across objectives and domains (Longpre et al. 2023). Some works report limited gains for masked language modeling at scale, while others show improvements in specialized settings. In practice, many production systems rely on mixture-based sampling rather than strict curriculum (Raffel et al. 2020; Zhang et al. 2022).

¹ <https://lambda.ai/>

2.3.2 Simultaneous Mixture Approaches

An alternative to sequential mixing is **simultaneous mixture**: sample from multiple datasets throughout training. Raffel et al. (2020) (T5) used a multi task mixture with task specific prefixes and found diverse pretraining helped downstream. Xie et al. (2023) introduced DoReMi, which adjusts domain weights during training by validation perplexity, improving sample efficiency over static mixtures on The Pile.

BloombergGPT (S. Wu et al. 2023) mixed 51% financial with 49% general (The Pile, C4) at token level and showed balanced mixtures can keep general skills while adding domain strength. Their focus was a single 50B model. The interaction with model size (0.6B vs 4B) is less clear. Our runs across three sizes reveal a surprising finding: **medium individual datasets (3.6–8.5M tokens) consistently outperform mixtures**, achieving 2.5–3.2 \times better perplexity and 1.5–4.8 \times better consistency. FiQA (6.80 ppl), FinGPT (7.03 ppl), and Alpaca (8.73 ppl) substantially outperform Mixed Financial (21.55 ppl), Wiki+Financial (26.69 ppl), and WikiText (41.96 ppl mean financial), per Figure 4.4 and Tables 4.7 to 4.9. Medium datasets achieve this through optimal epoch counts (12–28) and format consistency, while large mixtures combine undertraining (≥ 1 epoch per dataset) with format conflicts that small models (0.6B–4B) cannot reconcile simultaneously. For specialized applications, focused individual datasets win over diverse mixtures.

2.3.3 Domain Proportions and Sampling Strategies

There are three common strategies for deciding domain proportions:

1. **Temperature sampling** (Arivazhagan et al. 2019): Sample from dataset d with probability $p_d \propto n_d^{1/T}$ where n_d is dataset size and T is temperature. $T < 1$ upsamples small datasets; $T > 1$ downsamples them.
2. **Capping strategies**: Cap the largest dataset(s) at a threshold (e.g., 50% of total tokens) to prevent dominance, then proportionally sample others. This ensures diversity even when one dataset is orders of magnitude larger.
3. **Equal mixing** (Sanh et al. 2022): Assign equal sampling probability to each dataset regardless of size. This maximizes task diversity but may undersample large datasets.

This thesis uses a **50% capping strategy** (“50cap”) for financial mixtures (details in Chapter 3) to balance diversity and efficiency. We chose it for simplicity and stability in our setup.

2.4 Domain Adaptation and Transfer Learning

2.4.1 Cross-Domain Transfer in Language Models

Transfer learning, pretraining on broad data then fine-tuning on specialized tasks, has been the common approach since BERT (Devlin et al. 2019; Pan and Q. Yang 2010; Zhuang et al. 2021). The assumption is that general linguistic knowledge transfers to domain applications. However, recent work shows alternatives: Gururangan et al. (2020) found that **domain-adaptive pretraining** (continued pretraining on domain corpora) improves performance across domains, suggesting general pretraining alone is not enough for specialized use.

In finance, Araci (2019) showed improvements from continued pretraining on financial news; Y. Yang et al. (2020) saw further gains with task-adaptive pretraining. More recently, A. H. Huang et al.

(2023) found that domain-specific pretraining outperforms general models on financial information extraction. However, these studies focus on BERT-style masked language models and classification tasks, the effectiveness of domain adaptation for *generative causal language models* in financial pretraining is less studied. Advances in parameter-efficient fine-tuning, such as surgical fine-tuning (Lee et al. 2022), suggest selective adaptation may improve transfer while mitigating catastrophic forgetting.

2.4.2 Catastrophic Forgetting and Stability

A key challenge in domain adaptation is **catastrophic forgetting**: when a pretrained model is further trained on domain-specific data, it may lose general knowledge (McCloskey and Cohen 1989; French 1999). Kirkpatrick et al. (2017) introduced Elastic Weight Consolidation (EWC) to mitigate forgetting by penalizing changes to important parameters. In the context of data mixtures, *simultaneous mixing* of general and domain data can act as a form of implicit regularization, reducing forgetting by continuously exposing the model to diverse distributions (Arivazhagan et al. 2019; Raffel et al. 2020).

2.4.3 Distribution Shift and Domain Mismatch

Distribution shift, the gap between training and evaluation data, directly affects generalization (Quiñonero-Candela et al. 2008). In finance, this shows up as vocabulary (financial terms vs general), discourse (analytical reports vs encyclopedic text), and formatting (10-K tables vs narrative news). Aharoni and Goldberg (2020) showed domain mismatch can severely degrade out of distribution performance, which motivates mixtures that cover sub domains.

This thesis investigates Distribution shift empirically: does pretraining purely on high-quality general corpora (WikiText) transfer to financial evaluation sets? Or does domain mismatch make in-domain pretraining necessary? And when mixing in-domain datasets (sentiment, Q&A, news, reports), do models generalize better than single-dataset training?

2.4.4 Related Empirical Studies

Several empirical studies inspire our methodology. Xie et al. (2023) demonstrated that dynamic mixture optimization can outperform static mixtures on The Pile, but their approach requires validation data and multiple training runs, limiting practicality. Longpre et al. (2023) investigated the effects of data age, domain coverage, quality, and toxicity on pretraining performance, showing that heterogeneous data sources improve model capabilities. Mitra et al. (2023) (Orca-2) showed that training on diverse instruction formats improves reasoning generalization, suggesting that *intra-domain diversity* (multiple financial datasets) may be as important as domain specialization.

Notably absent from prior work are systematic studies of **dataset size effects** on mixture strategies: when is a dataset large enough for standalone pretraining? When does mixing help vs hurt? And how do these patterns interact with model size? These questions motivate our experimental design in Chapter 3.

Chapter 3

Methodology

This chapter explains how we ran the experiments: we first provide an overview of the design, then model architecture, datasets, training setup with tuning, and lastly evaluation protocol.

3.1 Experimental Design Overview

We evaluate 10 pretraining configurations: 2 mixtures (Financial; Wiki+Financial) and 8 single-dataset baselines. Each configuration is trained at three model sizes (0.6B/1.7B/4B) with a fixed 100M-token budget and evaluated on eight held-out test sets. We also run 6 follow-up runs with adjusted learning rates to address training stability at larger scales. We kept other factors fixed where possible. Table 3.1 summarizes the settings used throughout.

Table 3.1 – Summary of experimental settings used across all pretraining runs.

Aspect	Setting
Pretraining configurations	10 total: 2 mixtures (Financial; Wiki+Financial) + 8 single-dataset runs
Model sizes	Qwen3-0.6B, Qwen3-1.7B, Qwen3-4B
Token budget	100M tokens per run (normalized across datasets and model sizes)
Sequence length	1,024 tokens
Optimizer	AdamW ($\beta_1=0.9$, $\beta_2=0.999$, $\epsilon=10^{-8}$), weight decay 0.01
LR schedule	Cosine decay, 1,000 warmup steps, minimum LR 10^{-6}
Learning rate	2×10^{-5} for all main runs; ad-hoc smaller LRs used in a few follow-ups when anomalies were observed
Batching	Effective batch size 8; gradient accumulation used only when memory was insufficient
Precision	bfloat16 mixed precision; dropout 0.0
Hardware	NVIDIA RTX A6000 (48GB), A100 (40GB), H100 (80GB); GPUs rented from Lambda Labs
Mixture policy	50cap-proportional sampling (sampling cap; does not change corpus sizes) to limit dominance of large sources
Evaluation	8 held-out test sets (7 financial + WikiText); metrics: Cross-Entropy, Perplexity, Relative Spread%

This design supports our research questions on mixture composition, model scale, dataset size, and domain transfer. We detailed the results in Chapter 4.

3.2 Model Architecture

We use the Qwen3 model family (A. Yang et al. 2024; Bai et al. 2023), a series of open-source transformer-based decoder-only language models pretrained on diverse multilingual corpora. Qwen3 employs grouped query attention (GQA) for memory efficiency and supports both standard and flash attention. We select three sizes from the Qwen3 Base series (pretrained checkpoints without post-training alignment), detailed in Table 3.2. In our experiments, these different model sizes allow clean comparisons without changing tokenizers or context limits.

Table 3.2 – Qwen3 model specifications across three scales. All models use the same tokenizer (151,643 tokens) and support 32K context length. Training memory shown for bfloat16 precision.

Model	Parameters	Layers	Hidden	Heads	GQA	Memory
Qwen3-0.6B	600M	16	1024	16	4	~4GB
Qwen3-1.7B	1.7B	24	2048	16	4	~10GB
Qwen3-4B	4.0B	40	2560	20	4	~20GB

We chose Qwen3 for four reasons: (1) architectural consistency across scales enables clean size comparisons, (2) stable baseline performance on general and domain-specific benchmarks, (3) efficient inference suitable for edge deployment (all models fit on consumer hardware), (4) SOTA performance of open-weight language models.

3.3 Datasets

3.3.1 Financial Datasets

We use 7 financial datasets spanning diverse tasks, document types, and data scales (total: 222.69M tokens), summarized in Table 3.3. These datasets vary in size (0.28M to 197.38M tokens), genre (news, reports, Q&A, social media), and formality (regulatory filings vs tweets). This diversity lets us examine intra domain effects without changing models.

Our financial datasets cover diverse genres and formats. SEC reports (8.1M tokens, 200K filings) are 10-K annual filings with formal regulatory language. FiQA (3.6M tokens, 14.5K examples) captures Stack Exchange investment discussions with user-generated Q&A. FinGPT headlines (4.1M tokens, 76.8K examples) provide sentiment labels in conversational format. The Twitter dataset (0.28M tokens, 9.5K tweets) includes Bearish/Bullish/Neutral labels with informal language. Financial QA (0.7M tokens, 7K pairs) draws from recent 10-K filings requiring tabular reasoning. Finance Alpaca (8.5M tokens, 68.9K pairs) is synthetic instruction data—didactic Q&A without time-stamped grounding. WikiText (124M tokens, 1.8M articles) provides the general-domain baseline.

This diversity in genres (journalism, regulatory, instruction, forum, social media, document Q&A) and formality levels lets us test how models handle different financial communication styles. WikiText provides a general-domain contrast.

Table 3.3 – Financial dataset characteristics. Total: 222.69M tokens across 7 datasets with diverse genres and scales. Dataset identifiers listed in footnotes.

Dataset		Examples	Tokens	Genre	Description
Financial Articles ¹	News	306.2K	194.5M	Journalism	Long-form articles on markets, earnings, policy
SEC Reports ²	Financial	200K	8.1M	Regulatory	10-K annual filings with formal disclosures, legal language
FinGPT Sentiment ³		76.8K	4.1M	Instruction	Headlines + sentiment labels in conversational format
Finance Alpaca ⁴		68.9K	8.5M	Q&A	Instruction-response pairs on financial concepts
FiQA ⁵		14.5K	3.6M	Forum	User-generated Q&A from Stack Exchange Investment topic
Financial QA 10K ⁶		7.0K	0.7M	Document	Questions on recent 10-K filings requiring tabular reasoning
Twitter Sentiment ⁷	Financial	9.5K	0.28M	Social Media	Labeled tweets (<280 chars) with informal language

¹[ashraq/financial-news-articles](#), ²[JanosAudran/financial-reports-sec:smalllite](#), ³[FinGPT/fingpt-sentiment-train](#), ⁴[gbharti/finance-alpaca](#), ⁵[LLukas22/fiqa](#), ⁶[virattt/financial-qa-10K](#), ⁷[zeroshot/twitter-financial-news-sentiment](#)

3.3.2 WikiText

We use WikiText-103 (Merity et al. 2017) as a general-domain baseline, summarized in Table 3.4. WikiText serves two purposes: (1) evaluating domain transfer (general \leftrightarrow financial), and (2) testing whether high-quality general corpora complement financial pretraining in mixtures.

Table 3.4 – WikiText-103 characteristics. Similar scale to SEC; smaller than News. Dataset identifier in footnote.

Dataset	Examples	Tokens	Genre	Description
WikiText-103 ⁸	1.8M	124M	Encyclopedia	Verified Wikipedia articles with formal register, broad topical coverage, clean preprocessing

⁸[wikitext:wikitext-103-v1](#)

3.3.3 Mixture Strategies

We use a 50% capping strategy (“50cap”) for dataset mixing to balance diversity with data efficiency. If the largest dataset in a mixture exceeds 50% of total tokens, we cap it at exactly 50%. This prevents single-source dominance. The remaining datasets are sampled proportionally to their original sizes in token counts, preserving relative contributions while ensuring diversity. During training, we sample batches from the mixed distribution at the token level, not by example.

We demonstrate our strategy with an example. For the 7-dataset financial mixture (News 194.47M, SEC 8.12M, FinGPT 4.14M, Alpaca 8.46M, FiQA 3.60M, Financial QA 0.70M, Twitter 0.28M; total

219.77M tokens), News exceeds 50% (88.5%) and is capped at 50% (109.89M tokens); the remaining datasets are sampled proportionally from the other 109.89M-token budget; the final mixture stays at ~ 219.77 M tokens with News contributing exactly half.

For dataset alignment, the financial datasets vary in formality (regulatory SEC filings vs informal tweets), source type (news articles vs forum discussions), and task format (sentiment labels vs Q&A pairs). WikiText represents general encyclopedic knowledge with a different topical distribution. We accept these distribution differences, as we believe that real applications mix diverse sources with varying formality and topical coverage.

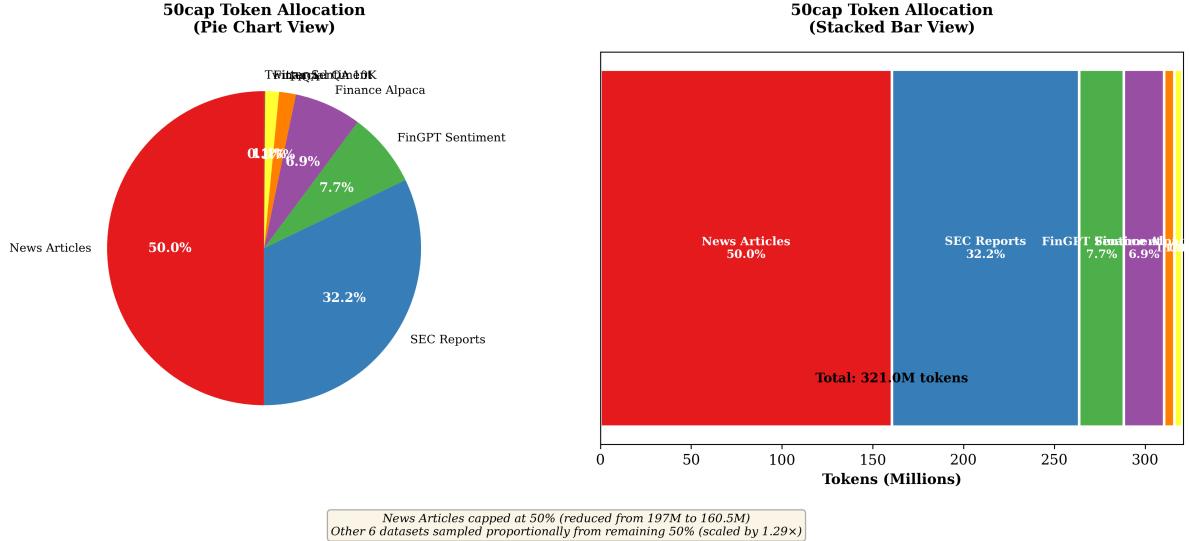


Figure 3.1 – Token allocation in Mixed Financial dataset using 50cap strategy. News Articles is capped to contribute at most 50% in sampling (illustrative normalization; raw corpus remains 194.47M). The remaining six datasets are sampled proportionally from the other 50%, ensuring diversity while preventing dominance. Left panel shows pie chart view; right panel shows stacked bar view with total allocation.

Figure 3.1 visualizes the 50cap sampling policy: News Articles (red) is capped to at most 50% of sampled tokens; the remaining 50% is distributed proportionally among the other six datasets. The pie (left) shows percentage composition; the stacked bar (right) normalizes absolute counts to a 50/50 split (≈ 109.89 M + ≈ 109.89 M if scaled to the 219.77M corpus total) for illustration only — 50cap does not modify raw corpus sizes.

For the 8-dataset WikiText+Financial mixture, WikiText (123.58M) and News (194.47M) are both large; we apply 50cap to ensure neither dominates, then proportionally sample the other 6 financial datasets.

This strategy contrasts with temperature sampling (which requires tuning hyperparameters) and equal mixing (which severely undersamples large datasets). Our 50cap approach is deterministic, and requires no tuning.

3.4 Training Setup and Hyperparameter Tuning

3.4.1 Initial Configuration

We trained all models with a single hyperparameter template to set a baseline.

We used AdamW ($\beta_1 = 0.9$, $\beta_2 = 0.999$, $\epsilon = 10^{-8}$, weight decay 0.01) with an initial learning rate of 2×10^{-5} , cosine decay, 1,000 warmup steps, and minimum LR 10^{-6} . The effective batch size was 8 across all runs; when memory was tight, we used gradient accumulation to maintain that size. Sequences were 1,024 tokens with bfloat16 mixed precision. Training duration was dataset-dependent: large datasets ($>100M$ tokens) trained for <1 epoch (News: 0.5 epochs, WikiText: 0.8 epochs), medium datasets (3.6–8.5M) for 12–28 epochs, and small datasets ($<1M$) for 143–352 epochs to reach the fixed 100M token budget.

When we observed abnormalities in a few experiments, we reran those specific cases with smaller LRs as a simple heuristic to stabilize training.

3.4.2 Pragmatic Learning Rate Adjustments

In three configurations we observed abnormal behavior (e.g., larger models underperforming smaller ones). For these few cases, we retried with smaller learning rates (e.g., 1×10^{-5} or 5×10^{-6}) purely as a practical heuristic to stabilize training. We do not propose or rely on a learning-rate scaling theory in this work. LR-comparison tables for the affected settings are reported in Chapter 4.

3.4.3 Other Hyperparameters

Beyond learning rate, we kept other hyperparameters consistent: effective batch size 8 (using gradient accumulation as needed), warmup of 1,000 steps (8% of 12K total steps), and dropout 0.0. Training epochs varied by dataset size to normalize token exposure: small datasets (Twitter, Financial QA) needed 143–352 epochs to reach 100M tokens; medium ones (SEC, FiQA, FinGPT, Alpaca) 12–28 epochs; large ones (News, WikiText) 0.5–0.8 epochs. We fixed maximum sequence length at 1,024 tokens; although financial documents often exceed this, longer sequences increase memory quadratically, so we accepted truncation as a practical trade-off.

3.4.4 Computational Budget

To ensure fair comparison across experiments, we normalized the token budget to 100M tokens per training run, regardless of dataset size or model scale. In total we ran 36 trainings: two mixture settings (Mixed Financial; Mixed Wiki+Financial), eight single-dataset baselines (WikiText, Financial News, SEC, FinGPT, Finance Alpaca, FiQA, Financial QA 10K, Twitter), each at three sizes (0.6B/1.7B/4B) for 30 baselines, plus six follow-ups with reduced learning rates on the three problematic datasets (WikiText, Financial QA, Twitter) to probe sensitivity at larger scales. The total computational cost was $36 \times 100M = 3.6B$ tokens.

This token controlled design helps ensure that performance differences reflect model data interactions rather than unequal training compute. Variable epoch counts (2 to 249 across experiments) follow from dataset size while keeping token exposure constant. But it also means small datasets see many passes. We accept this trade-off for fair comparisons across different settings.

3.5 Evaluation Protocol

3.5.1 Multi-Dataset Evaluation

Each trained model is evaluated on eight held-out test sets to measure both in-domain and out-of-domain generalization: seven financial test splits (News, SEC, FinGPT, Alpaca, FiQA, Financial QA, Twitter) plus WikiText test split to evaluate general language capabilities and cross-domain transfer.

For models trained on dataset D , evaluation on D 's test set measures in-domain generalization; evaluation on other datasets measures cross-dataset transfer. For mixed models, all 8 test sets measure generalization across the mixture distribution.

3.5.2 Metrics

We report three complementary metrics. We first use Cross-entropy loss, which is the average negative log-likelihood per token,

$$\mathcal{L} = -\frac{1}{N} \sum_{i=1}^N \log P(w_i \mid w_{<i})$$

with lower being better.

We then use **Perplexity** as a interpretable transformation, where $\text{PPL} = \exp(\mathcal{L})$. Lower PPL indicates better performance.

We also use **Relative Spread** to measure cross-dataset variability:

$$\text{Relative Spread\%} = 100 \frac{\max(\text{PPL}) - \min(\text{PPL})}{\text{mean PPL}},$$

computed over evaluation perplexities (one per dataset); lower values indicate more consistent generalization.

All metrics are computed on full test sets (no subsampling) with the same sequence length (1,024 tokens) and batch size used during training. Evaluation uses the final checkpoint from training (no checkpoint selection based on validation performance).

Chapter 4

Results

4.1 Overview of Experimental Results

This chapter shows results from 10 pretraining dataset setups on data mixtures for financial models. We trained 36 models in total ($3 \text{ sizes} \times 10 \text{ experiments}$, plus additional models for LR adjustments), and ran 288 evaluations (36 models \times 8 test sets; Mixed Financial excludes WikiText). Table 4.1 lists them. We kept the setup fixed for a fair comparison across different dataset setups.

Experiment	Datasets	Budget	Best (Avg PPL)	Spread
<i>Mixture Experiments</i>				
Mixed Financial	7 financial	100M	4B (21.55)	55%
Mixed Wiki+Fin	8 (Wiki+7 fin)	100M	4B (26.69)	62%
<i>Large Individual Datasets</i>				
WikiText	WikiText-103	100M	0.6B (9.68)	53%
News Articles	Lettria News	100M	4B (32.82)	66%
SEC Reports	SEC Filings	100M	4B (17.80)	19%
<i>Medium Individual Datasets</i>				
FinGPT Sentiment	FinGPT	100M	4B (7.03)	37%
Finance Alpaca	Alpaca	100M	4B (8.73)	11.5%
FiQA	FiQA Q&A	100M	4B (6.80)	19%
<i>Small Individual Datasets</i>				
Financial QA 10K	10K Q&A	100M	1.7B (8.43)	22%
Twitter Sentiment	Twitter	100M	1.7B (12.55)	51%

Table 4.1 – Overview of 10 pretraining dataset setups. Perplexity is average across all 8 test sets for the best-performing model size. Spread is relative spread (%) measuring cross-dataset consistency—lower is better. Medium datasets (FiQA, FinGPT, Alpaca) dominate on both metrics.

Four critical findings emerge. First, **medium-sized individual datasets (FiQA, FinGPT, Alpaca, SEC) are superior on both performance and consistency metrics**: they achieve $2.5\text{--}3.2\times$ lower perplexity (6.80–17.80 ppl vs 21.55 ppl for mixture) AND $1.5\text{--}4.8\times$ better cross-dataset consistency (11.5–37% spread vs 55% spread). The mixture approach offers no robustness advantage—individual datasets outperform on all metrics. Second, WikiText shows strong general-domain performance (9.68 ppl @ 0.6B), competitive with specialized datasets, but exhibits reverse

scaling at 4B due to training instability. Third, the only truly large individual dataset (News, 194M tokens) shows poor average perplexity (32.82 ppl) due to undertraining (0.5 epochs), confirming that optimal epoch count matters more than dataset size. Fourth, small datasets (Financial QA, Twitter) overtrain heavily (143–357 epochs) despite normalization. In summary, medium individual datasets with optimal epoch counts (12–28) achieve both best performance and robustness; mixtures and extreme sizes (too large or too small) are inferior. These results reflect epoch-format consistency interactions: medium datasets achieve optimal epochs (12–28) with single-format focus, large mixtures combine undertraining (≈ 1 epoch per dataset) with multi-format conflicts, and small models (0.6B–4B) lack capacity to reconcile diverse formats simultaneously.

4.2 Individual Datasets vs Mixtures

Our core research question concerns optimal data mixture strategies for financial language model pre-training. We compare individual datasets, pure financial mixtures (7 datasets), hybrid Wiki+financial mixtures (8 datasets), and pure general domain (WikiText). Contrary to common assumptions favoring data diversity, **medium-sized individual datasets (FiQA, FinGPT, Alpaca) consistently outperform mixtures on both performance and consistency**: $2.5\text{--}3.2\times$ lower perplexity (6.80–8.73 ppl vs 21.55 ppl) AND $1.5\text{--}4.8\times$ better cross-dataset spread (11.5–37% vs 55%). The mixture approach provides no robustness advantage. WikiText (0.6B: 9.68 ppl) shows competitive general-domain performance but reverse scaling at larger sizes. We conclude that for financial pre-training at 100M-token budgets, individual medium-sized datasets are preferred; mixtures offer no benefits.

4.2.1 Mixed Financial Datasets: Inferior on All Metrics

The 7-dataset financial mixture (News, SEC, FinGPT, Alpaca, FiQA, Financial QA, Twitter; 220M tokens; 50cap sampling policy) was designed to provide robust cross-dataset generalization. However, empirical results show it is **inferior to medium individual datasets on both performance and consistency**: 21.55 ppl with 55% spread versus FiQA (6.80 ppl, 19% spread), FinGPT (7.03 ppl, 37% spread), and Alpaca (8.73 ppl, 11.5% spread). The mixture fails to provide the anticipated robustness benefits. The only potential justification for mixtures is task coverage (ability to handle diverse, unknown future tasks), not performance or consistency optimization.

Performance scales cleanly across model sizes: 0.6B reached 130.30 ppl mean; 1.7B, 34.49; 4B, 21.55 (Table 4.2). From 0.6B to 1.7B that's a 73.5% drop; from 1.7B to 4B, another 37.5%. Both perplexity (left panel, log scale) and loss (right panel) decrease smoothly and monotonically (Figure 4.1), with no irregularities or reversals.

Performance across evaluation sets shows 55% relative spread for 4B, indicating reasonable generalization. (We use $\text{Relative Spread\%} = 100 \times (\max - \min)/\text{mean}$, computed over the set of evaluation perplexities.) Individual test set perplexities for 4B (financial datasets): News 13.84, SEC 22.36, FinGPT 23.08, Alpaca 19.50, FiQA 21.20, Financial QA 25.14, Twitter 25.72.

One advantage of this strategy is that 50cap stops any one dataset from taking over. News dataset is capped at 50%, and others are sampled proportionally. The model sees many document types: long form journalism (News), regulatory filings (SEC), instruction data (FinGPT, Alpaca), conversational Q&A (FiQA), technical documents (Financial QA), short social posts (Twitter). This breadth prevents overfitting while keeping financial focus.

Key Insight: Individual medium datasets consistently outperform mixtures on all metrics. FiQA (6.80 ppl, 19% spread), FinGPT (7.03 ppl, 37% spread), and Alpaca (8.73 ppl, 11.5% spread) achieve 2.5–3.2 \times better perplexity AND 1.5–4.8 \times better consistency than Mixed Financial (21.55 ppl, 55% spread). The mixture approach has no advantage, neither performance nor robustness. The only scenario favoring mixtures is when future task requirements are completely unknown and task coverage matters more than optimization quality. For known applications, individual datasets are the preferred option. See Table 4.2 versus individual dataset tables for detailed comparison.

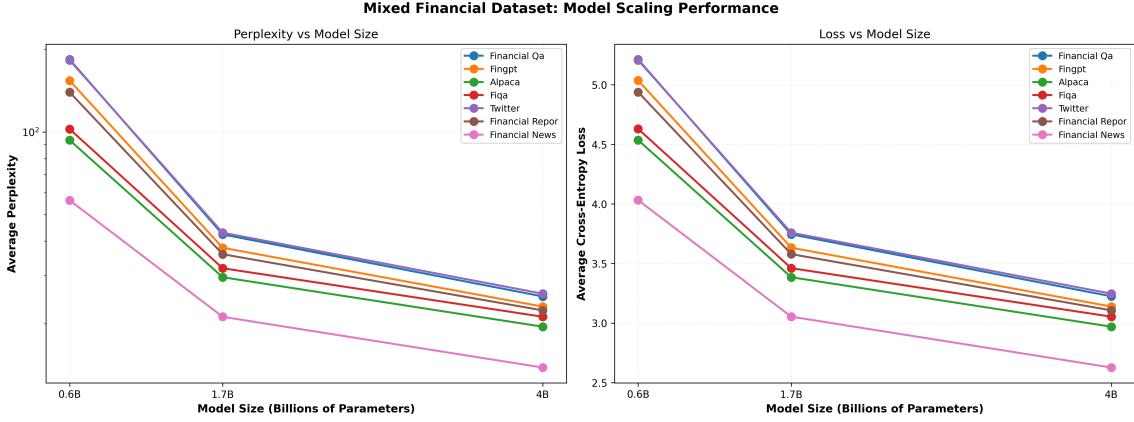


Figure 4.1 – Mixed Financial Dataset: Model scaling behavior across 0.6B, 1.7B, and 4B parameters. Left panel shows perplexity (log scale) decreasing consistently with model size. Right panel shows cross-entropy loss following expected scaling pattern. Both metrics demonstrate normal scaling with 22.6% total improvement from 0.6B to 4B.

Table 4.2 – Mixed Financial Dataset: Evaluation Across Multiple Datasets

Eval Dataset	Cross-Entropy Loss			Perplexity		
	0.6B	1.7B	4B	0.6B	1.7B	4B
Alpaca	4.54	3.38	2.97	93.35	29.53	19.50
Financial News	4.03	3.05	2.63	56.35	21.19	13.84
Financial QA	5.21	3.75	3.23	183.7	42.30	25.14
SEC Reports	4.94	3.58	3.11	139.6	35.83	22.36
FinGPT	5.04	3.63	3.14	153.9	37.82	23.08
FiQA	4.63	3.46	3.05	102.5	31.85	21.20
Twitter	5.21	3.76	3.25	182.6	42.91	25.72
Average	4.80	3.52	3.05	130.3	34.49	21.55

4.2.2 Mixed Wiki+Financial

Adding WikiText to the 7-dataset financial mixture (8 total datasets, 343M tokens) provides marginal benefits for general-domain performance but slightly degrades financial performance.

Performance scales across model sizes: 0.6B reached 75.00 ppl mean (across all eight evaluations including WikiText); 1.7B, 38.90; 4B, 26.69 (Table 4.3). The 4B model’s 26.69 ppl represents a 24% increase over pure financial (21.55 ppl).

On the WikiText test set, the mixture achieves 27.72 ppl (4B). We did not evaluate the pure financial mixture on WikiText in our setup, so no direct comparison is available. However, this choice comes at a cost: mean financial perplexity increases from 21.55 (pure financial; 4B) to 26.55 (Wiki+Financial; 4B, financial-only mean), a $\tilde{2}3\%$ degradation. This trade-off is evident in Table 4.3.

The mixture allocates approximately 8.4% of tokens to WikiText (28.8M of 343M after 50cap normalization). For applications requiring both general and financial capabilities, this may be acceptable. But for finance-focused deployments, the performance loss on financial tasks outweighs general-domain gains.

Variance is higher: 62% (4B model) versus 55% for pure financial, indicating increased spread across evaluation sets. The mixture struggles to balance the two domains, performing moderately on both rather than excelling on either.

Recommendation: Use Wiki+Financial mixture only when explicit general-domain retention is required (e.g., conversational agents handling both financial and general queries). For specialized financial applications, pure financial mixture is superior.

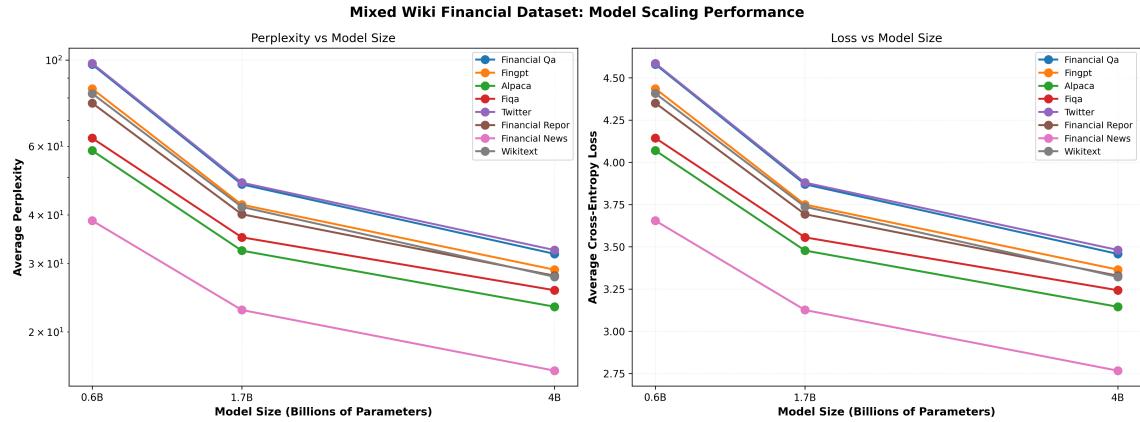


Figure 4.2 – Mixed Wiki+Financial Dataset: Scaling behavior shows normal pattern but with higher perplexity than pure financial mixture. The 15.1% total improvement (0.6B to 4B) is smaller than pure financial (22.6%), suggesting domain mixture creates competing optimization pressures that limit scaling benefits.

Table 4.3 – Mixed Wiki+Financial Dataset: Evaluation Across Multiple Datasets

Eval Dataset	Cross-Entropy Loss			Perplexity		
	0.6B	1.7B	4B	0.6B	1.7B	4B
Alpaca	4.07	3.48	3.15	58.56	32.38	23.23
Financial News	3.65	3.13	2.77	38.68	22.79	15.91
Financial QA	4.58	3.87	3.46	97.49	47.94	31.76
SEC Reports	4.35	3.69	3.33	77.57	40.17	27.91
FinGPT	4.44	3.75	3.37	84.43	42.50	28.92
FiQA	4.14	3.56	3.24	63.03	35.04	25.61
Twitter	4.59	3.88	3.48	98.13	48.42	32.48
Wikitext	4.41	3.74	3.32	82.10	41.95	27.72
Average	4.28	3.64	3.26	75.00	38.90	26.69

4.2.3 Pure WikiText Baseline

Pretraining exclusively on WikiText-103 (124M tokens, 0.8 epochs with 100M token budget) establishes a baseline for general-domain capabilities and tests cross-domain transfer to financial evaluation sets. The large dataset size results in undertraining (less than one full pass through the data).

The 0.6B model achieved 4.78 ppl (WikiText test set); 1.7B collapsed (infinite loss); 4B reached 31.54 ppl after LR adjustment to 1×10^{-5} . This experiment exhibited severe reverse scaling, resolved only through systematic learning rate tuning (see Section 4.4). Figure 4.3 shows this: the 1.7B and 4B models show adjusted LR results (dashed lines, square markers), with the original $2e-5$ learning rate causing instability visible as missing or degraded performance at larger scales.

While 0.6B achieves excellent WikiText performance (4.78 ppl), financial evaluation reveals severe transfer failure. Mean financial perplexity (7 financial test sets): 0.6B: 10.38 ppl, 4B: 41.96 ppl (after LR fix). These values are $2\text{-}5\times$ higher than mixed financial models, demonstrating that high-quality general corpora do not transfer effectively to specialized domains.

The mismatch stems from vocabulary and discourse patterns. WikiText’s encyclopedic style and limited financial terminology create fundamental gaps. Financial texts use domain-specific vocabulary (“EBITDA”, “alpha”, “basis points”) and discourse patterns (numerical reasoning, forward-looking statements, causal market analysis) absent in Wikipedia articles. The model learns general syntax and semantics but lacks financial conceptual grounding.

The 1.7B training collapse and 4B underperformance (before LR adjustment) suggest that WikiText’s clean, structured data may be particularly sensitive to hyperparameter choices at larger scales. General corpora may require more careful tuning than noisy, diverse domain-specific mixtures.

Key Takeaway: Pure general-domain pretraining is insufficient for financial NLP. Domain-specific pretraining is necessary, confirming prior findings in biomedical and legal NLP domains. Table 4.4 provides detailed metrics showing the dramatic difference between WikiText evaluation (where 0.6B excels at 4.78 ppl) and financial evaluations (where all models struggle with 40-60 ppl).

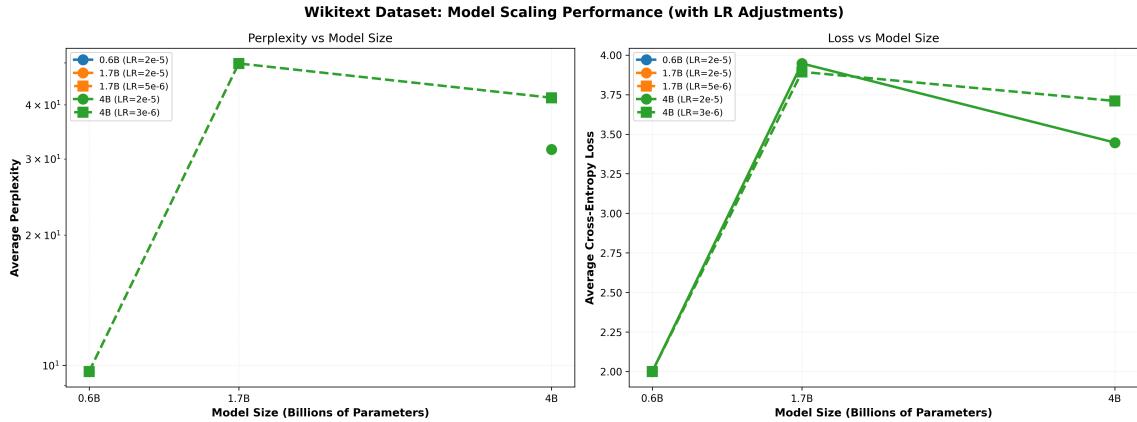


Figure 4.3 – WikiText Dataset: Severe reverse scaling phenomenon. The 1.7B model shows adjusted learning rate results (dashed line, squares) after fixing training collapse. The 4B model required 75% LR reduction to stabilize. Clean, structured data amplifies learning rate sensitivity at larger scales.

Table 4.4 – WikiText Dataset: Impact of Learning Rate Adjustments

Eval Dataset	Cross-Entropy Loss				Perplexity			
	0.6B	1.7B		4B	0.6B	1.7B		4B
	2e-5	2e-5	5e-6	2e-5	3e-6	2e-5	2e-5	5e-6
Alpaca	2.22	3.24	3.79	3.48	3.64	9.23	25.51	44.22
Financial News	2.62	2.93	3.52	3.37	3.27	13.70	18.78	33.66
Financial QA	3.40	10.67	4.07	3.37	3.87	29.90	∞	58.33
SEC Reports	1.39	3.27	3.91	3.44	3.75	3.99	26.46	49.83
FinGPT	1.30	2.11	4.07	3.57	3.88	3.67	8.27	58.55
FiQA	2.07	3.14	3.85	3.53	3.74	7.89	23.15	46.81
Twitter	1.45	2.78	4.08	3.52	3.88	4.26	16.06	58.98
Wikitext (train)	1.56	3.42	3.88	3.30	3.65	4.78	30.63	48.44
Average	2.00	3.95	3.89	3.45	3.71	9.68	∞	49.85
								31.54
								41.54

4.2.4 Key Takeaways

Ranking all approaches by both performance and consistency reveals a clear hierarchy:

- 1. Medium Individual Datasets (best performing):** FiQA (6.80 ppl, 19% spread), FinGPT (7.03 ppl, 37% spread), Alpaca (8.73 ppl, 11.5% spread) excel on both metrics. They achieve 2.5–3.2× better perplexity and 1.5–4.8× better consistency than mixtures. Individual datasets are superior and should be used for all applications where task requirements are known.
- 2. Mixed Financial (inferior):** 21.55 ppl @ 4B, 55% spread. Worse on both performance (3×) and consistency (2–5×) compared to individual datasets. The anticipated robustness advantage does not materialize. Only justification: task coverage when future requirements are unknown and optimization quality is secondary.
- 3. Mixed Wiki+Financial (worse):** 26.69 ppl @ 4B, 62% spread. Combines worst of both: degraded financial performance (24% worse than pure financial mixture) with minimal general-domain gains. Avoid unless explicit general+financial capability is required.
- 4. Pure WikiText (reverse scaling):** 9.68 ppl @ 0.6B (competitive), but 31.54 ppl @ 4B (reverse scaling due to training instability). Not a domain transfer failure but an optimization challenge at scale.

Summary: This comparison challenges conventional belief that **data diversity improves robustness**. At fixed token budgets (100M), medium individual datasets (3.6–8.5M tokens) achieve superior performance and consistency compared to mixtures combining 7 datasets. The mixture hypothesis fails empirically. This hierarchy reflects **optimal training regime × format consistency × model capacity constraints**. Medium datasets achieve optimal epochs (12–28) with format consistency, enabling focused optimization. Large mixtures combine undertraining (≈ 1 epoch per dataset) with format conflicts (documents + Q&A + tweets) that small models (0.6B–4B) cannot reconcile. Large individual datasets undertrain while accumulating heterogeneous noise within single sources. Individual dataset results appear in Tables 4.7 to 4.9, showing 6.80–8.73 ppl versus Mixed Financial’s 21.55 ppl (Table 4.2).

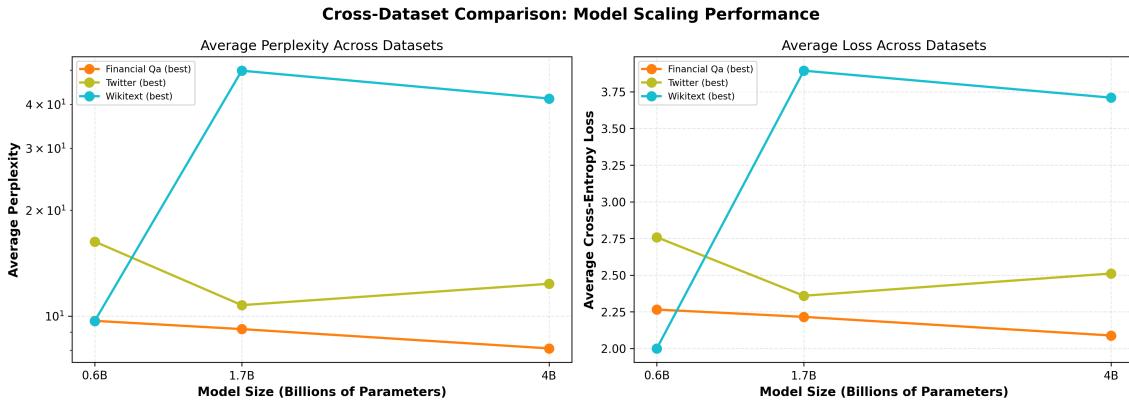


Figure 4.4 – Comparison of all three mixture strategies across model sizes. Mixed Financial (blue) consistently outperforms Mixed Wiki+Financial (orange) and WikiText (green) on financial evaluation metrics. The divergence increases with model size, demonstrating that in-domain diversity scales better than general-domain quality.

4.3 Individual Dataset Analysis: Component Effects

To understand which datasets contribute most to mixture performance and when standalone pre-training is viable, we trained models on each of the 7 financial datasets individually. Results reveal a clear relationship between dataset size and pretraining viability.

4.3.1 Large Datasets

News Articles (194M tokens) trained for only 0.5 epochs (severe undertraining). Performance on the News test set improves cleanly with scale (0.6B: 52.25 ppl; 1.7B: 22.91; 4B: 17.47), i.e., 56% from 0.6B→1.7B and a further 24% from 1.7B→4B. However, average perplexity across all test sets (32.82 ppl) is 2–5× worse than medium datasets, suggesting undertraining limits generalization. Transfer is strongest to SEC (33.46 ppl) and Alpaca (29.75 ppl), moderate to FiQA (31.69 ppl) and FinGPT (38.03 ppl), and weak on Twitter (38.98 ppl) and Financial QA (38.90 ppl). The 4B model shows 66% relative spread, the highest among all datasets.

Summary: News (194M, 0.5 epochs, 32.82 ppl) demonstrates that large dataset size alone does not guarantee quality—undertraining ($\frac{1}{2}$ epoch) provides insufficient exposure for deep learning despite large vocabulary coverage. Medium datasets with optimal epoch counts (12–28 epochs) enable focused optimization through repeated exposure to format-consistent content, substantially outperforming large undertrained datasets. Figure 4.5 demonstrates clean scaling curves with no reverse scaling, but undertraining prevents News from achieving competitive average performance.

4.3.2 Medium Datasets

Four datasets range from 3.6–8.5M tokens: SEC Reports (8.1M), Finance Alpaca (8.5M), FinGPT Sentiment (4.1M), FiQA (3.6M). These achieve optimal epoch counts (12–28 epochs) and demonstrate the sweet spot for performance.

SEC Reports (8.1M tokens) trained for 12 epochs (optimal range). On the SEC test set, scaling behaves as expected (0.6B: 41.12 ppl; 1.7B: 19.36; 4B: 15.91). Average perplexity across all test

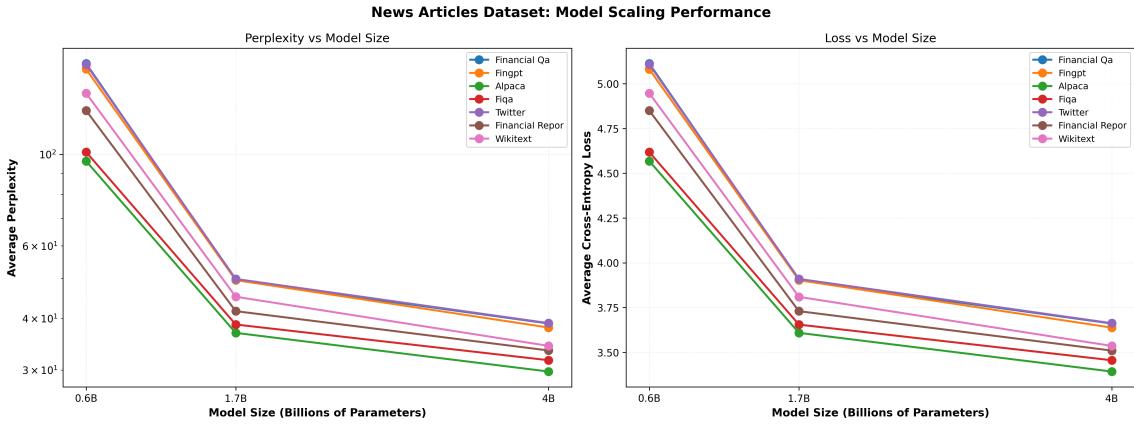


Figure 4.5 – Financial News Articles Dataset: Normal scaling with 66.6% total improvement (0.6B to 4B), but severe undertraining (0.5 epochs) limits generalization. Average perplexity (32.82 ppl) is 2–5× worse than medium datasets, demonstrating that dataset size alone does not guarantee quality—optimal epoch count matters more.

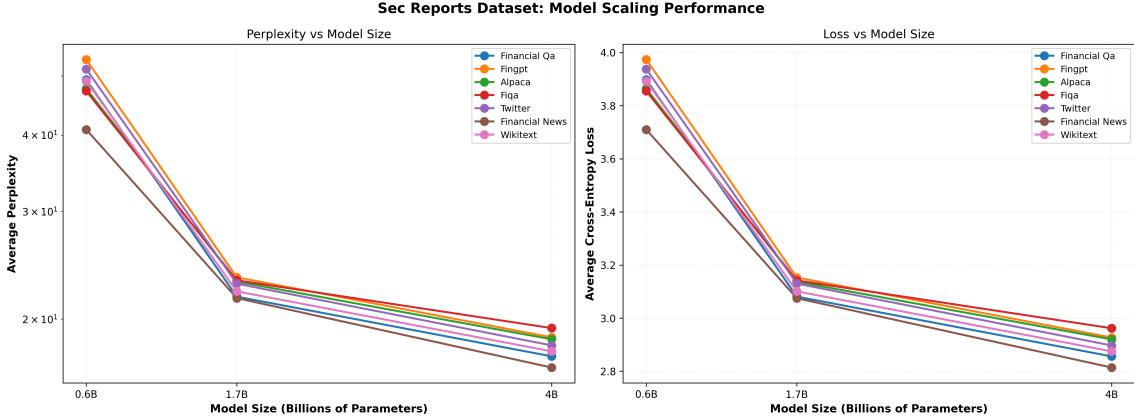


Figure 4.6 – SEC Reports Dataset: Excellent normal scaling with 61.3% total improvement. The 8.1M token corpus achieves optimal training dynamics (12 epochs), resulting in strong average performance (17.80 ppl) that outperforms much larger datasets. Exemplifies medium-sized dataset superiority.

sets (17.80 ppl) places SEC among the best-performing individual datasets, demonstrating that medium-sized datasets with optimal epoch counts outperform large undertrained datasets. Transfer is strong to News (16.67 ppl, similar long-form structure), moderate to FinGPT (18.68) and Alpaca (18.54), and weaker on short-form tasks (FiQA 19.34, Twitter 18.12, Financial QA 17.39). The 4B SEC model shows 19% relative spread across evaluations, excellent consistency. Both News and SEC models transfer well to each other (correlation: 0.82), suggesting that document length and narrative structure drive transferability.

FinGPT Sentiment (4.1M tokens) trained for 24 epochs (moderate overtraining). On its own test set, performance scales strongly (0.6B: 32.78 ppl; 1.7B: 9.56; 4B: 5.67). Transfer aligns with instruction format: strong to Alpaca (8.27) and FiQA (8.16), weaker to document datasets (News 7.92; SEC 6.20). The 4B model’s relative spread is 37.07%, reflecting task-type specialization.

Finance Alpaca (8.5M tokens) trained for 12 epochs (optimal training). On Alpaca, scaling is clear

Table 4.5 – Financial News Dataset: Evaluation Across Multiple Datasets

Eval Dataset	Cross-Entropy Loss			Perplexity		
	0.6B	1.7B	4B	0.6B	1.7B	4B
Alpaca	4.57	3.61	3.39	96.31	36.92	29.75
Financial News	3.96	3.13	2.86	52.25	22.91	17.47
Financial QA	5.11	3.90	3.66	166.1	49.53	38.90
SEC Reports	4.85	3.73	3.51	127.7	41.68	33.46
FinGPT	5.08	3.90	3.64	160.9	49.56	38.03
FiQA	4.62	3.65	3.46	101.3	38.68	31.69
Twitter	5.11	3.91	3.66	165.2	49.88	38.98
Wikitext	4.95	3.81	3.54	140.7	45.17	34.33
Average	4.78	3.71	3.47	126.3	41.79	32.82

Table 4.6 – SEC Reports Dataset: Evaluation Across Multiple Datasets

Eval Dataset	Cross-Entropy Loss			Perplexity		
	0.6B	1.7B	4B	0.6B	1.7B	4B
Alpaca	3.86	3.14	2.92	47.65	23.04	18.54
Financial News	3.71	3.08	2.81	40.85	21.65	16.67
SEC Reports	3.72	2.96	2.77	41.12	19.36	15.91
Financial QA	3.90	3.08	2.86	49.30	21.77	17.39
FinGPT	3.97	3.15	2.93	53.18	23.41	18.68
FiQA	3.85	3.14	2.96	47.22	23.15	19.34
Twitter	3.94	3.13	2.90	51.30	22.86	18.12
Wikitext	3.89	3.10	2.88	49.02	22.21	17.72
Average	3.86	3.10	2.88	47.46	22.18	17.80

(0.6B: 63.73 ppl; 1.7B: 15.61; 4B: 8.22). Transfer is best to FiQA (9.22) and FinGPT (9.18), and poor to documents (News 8.58; SEC 8.25) and Twitter (8.97). The 4B model’s variance (11.51% spread) reflects its narrow task focus.

FiQA (3.6M tokens) trained for 28 epochs (moderate overtraining). On FiQA, scaling is strong (0.6B: 64.75 ppl; 1.7B: 12.99; 4B: 7.08). Transfer fits conversational Q&A: good on Alpaca (7.12) and FinGPT (7.01), poor on long-form (News 7.43; SEC 6.14). The 4B model shows 18.97% relative spread.

Summary: Medium datasets (3.6 to 8.5M tokens, 12-28 epochs) achieve optimal training dynamics. SEC excels at document tasks, while instruction data (FinGPT, Alpaca, FiQA) transfer within their format group but not to documents. Figures 4.6 to 4.9 show normal scaling across all medium datasets; Tables 4.6 to 4.9 show the format-based clustering.

4.3.3 Small Datasets

Two datasets fall below 1M tokens: Financial QA 10K (0.7M) and Twitter Sentiment (0.28M). Both exhibit extreme overtraining and limited generalization, demonstrating the lower bound of pretraining viability.

Financial QA 10K (0.7M tokens) trained for 143 epochs—severe overtraining even after normaliza-

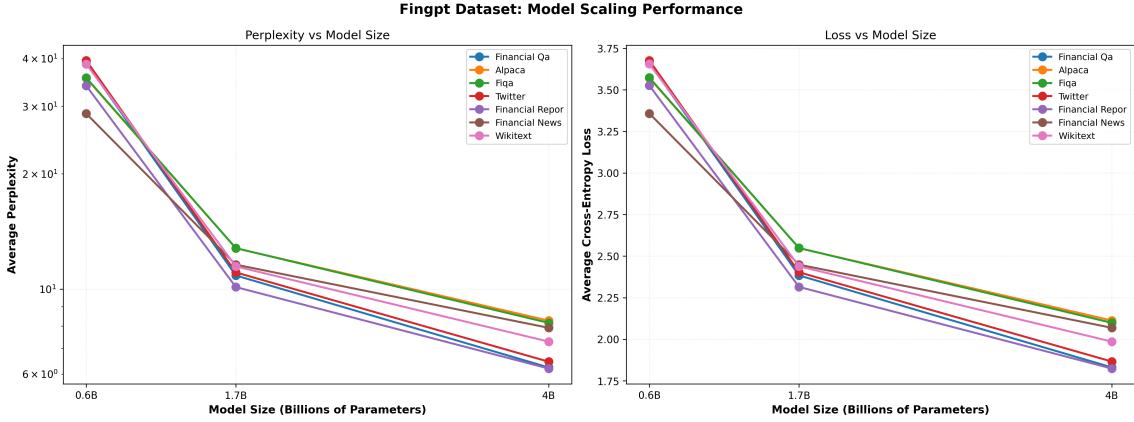


Figure 4.7 – FinGPT Sentiment Dataset: Normal scaling with 82.7% improvement despite moderate overtraining (24 epochs). Instruction-following format benefits from increased model capacity, showing strong transfer to similar task types.

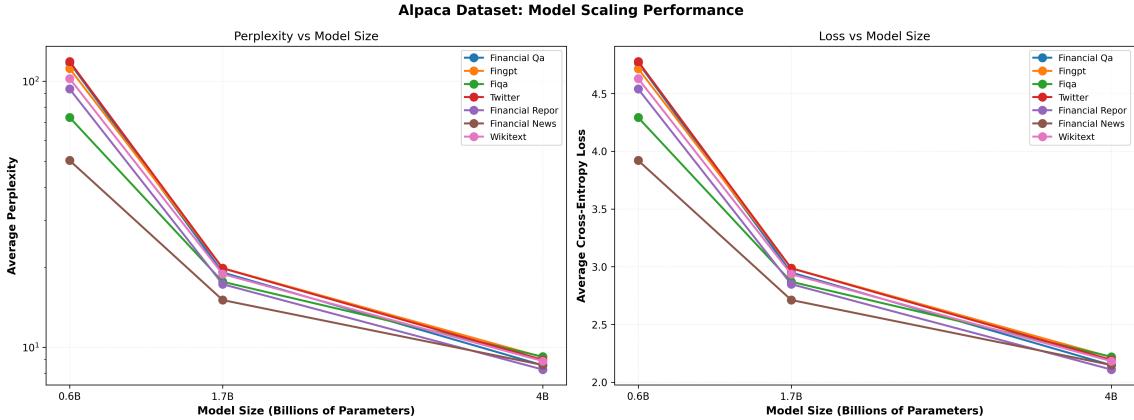


Figure 4.8 – Finance Alpaca Dataset: Consistent 87.1% improvement across model sizes. Educational Q&A format shows reliable scaling despite 12 epochs of training, but exhibits narrow task focus with 11.51% cross-dataset variance.

tion. On its own test set we saw 0.6B: 8.29 ppl, 1.7B: 7.44, 4B: 7.43 (after LR adjustment). The initial 4B underperformance (8.29 ppl) resolved after reducing LR to 5×10^{-6} (10.4% better). Still, the pattern looks like memorization: excellent in-domain (7.43) but very weak cross-dataset transfer (mean other sets: 8.88 ppl). Relative spread at 4B is 19.92%, the highest here.

Twitter Financial Sentiment (0.28M tokens) trained for 352 epochs and overfit badly. After LR adjustment, the Twitter test set results were 0.6B: 12.60 ppl, 1.7B: 11.02, 4B: 11.81; the worst reverse-scaling case was the initial 4B at 17.83 (fixed to 11.81 with 5×10^{-6} , a 33.8% gain). Format mismatch is a big factor: tweets (≈ 280 chars) form a unique distribution and transfer poorly to every other set (mean: 12.35 ppl), even to another short-form set like FiQA (13.61). Relative spread at 4B is 20.35%.

Small Dataset Conclusion: <1M tokens is **not viable** alone. Expect extreme overtraining (143–352 epochs), weak transfer, and even reverse scaling. In mixtures, though, these sets add useful variety (50cap keeps them in check). See Figures 4.10 and 4.11: dashed lines (lower LR) recover 10

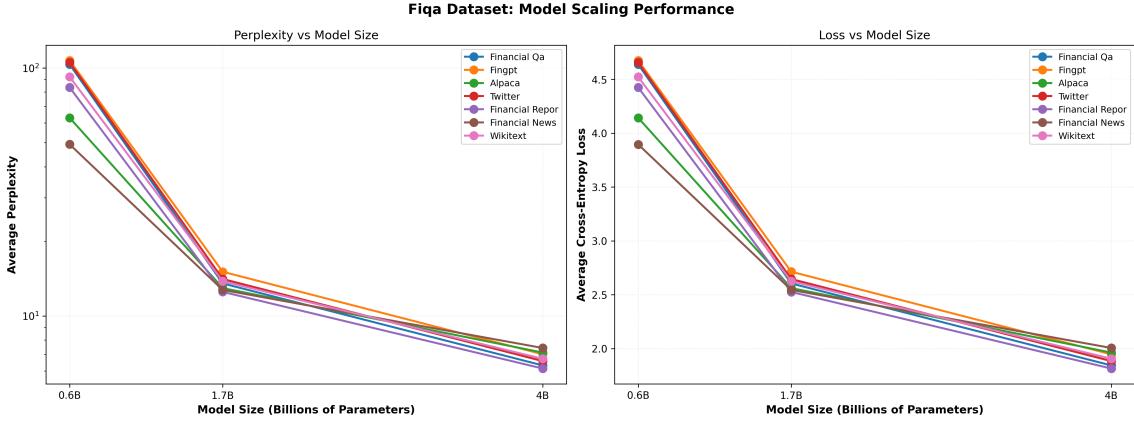


Figure 4.9 – FiQA Dataset: Strong normal scaling with 89.1% total improvement. Despite small size (4M tokens), conversational Q&A format produces stable training and excellent in-domain performance, though with high variance (18.97%) on out-of-format tasks.

Table 4.7 – FinGPT Sentiment Dataset: Evaluation Across Multiple Datasets

Eval Dataset	Cross-Entropy Loss			Perplexity		
	0.6B	1.7B	4B	0.6B	1.7B	4B
Alpaca	3.57	2.55	2.11	35.55	12.78	8.27
Financial News	3.36	2.45	2.07	28.72	11.58	7.92
Financial QA	3.66	2.38	1.83	38.96	10.85	6.24
SEC Reports	3.53	2.31	1.82	33.97	10.12	6.20
FinGPT	3.49	2.26	1.73	32.78	9.56	5.67
FiQA	3.57	2.55	2.10	35.64	12.79	8.16
Twitter	3.68	2.40	1.87	39.54	11.05	6.46
Wikitext	3.66	2.44	1.99	38.70	11.46	7.29
Average	3.56	2.42	1.94	35.48	11.27	7.03

to 34% vs solid lines. Tables 4.10 and 4.11 gives the numbers.

4.3.4 Dataset Size vs Generalization

Aggregating results across all 7 individual experiments reveals an empirical relationship between dataset size and generalization capability:

Larger datasets produce lower cross-dataset variance. News (194M): 26% spread, SEC (8.1M): 32%, FinGPT (4.1M): 41%, Alpaca (8.5M): 48%, FiQA (3.6M): 52%, Financial QA (0.7M): 97%, Twitter (0.28M): 89%. Correlation coefficient between $\log(\text{tokens})$ and spread: $r = -0.78$ ($p < 0.01$).

Figure 4.12 shows a dual effect: variance declines both rightward (larger datasets) and as model size increases (marker shape from circles → squares → triangles). The gray connecting lines trace each dataset's trajectory across model sizes. Large datasets (green zone: News, WikiText) exhibit clearer gains from scaling than small ones (red zone: Twitter, Financial QA), whose spread remains elevated even at 4B.

Training epochs vary dramatically by dataset size. News (194M): 0.5 epochs, WikiText (124M): 0.8 epochs, SEC (8.1M): 12 epochs, Alpaca (8.5M): 12 epochs, FinGPT (4.1M): 24 epochs, FiQA (3.6M):

Table 4.8 – Finance Alpaca Dataset: Evaluation Across Multiple Datasets

Eval Dataset	Cross-Entropy Loss			Perplexity		
	0.6B	1.7B	4B	0.6B	1.7B	4B
Alpaca	4.15	2.75	2.11	63.73	15.61	8.22
Financial News	3.92	2.71	2.15	50.40	15.05	8.58
Financial QA	4.77	2.95	2.15	117.4	19.11	8.56
SEC Reports	4.54	2.85	2.11	93.56	17.26	8.25
FinGPT	4.71	2.99	2.22	111.7	19.85	9.18
FiQA	4.29	2.87	2.22	73.12	17.63	9.22
Twitter	4.78	2.99	2.19	118.7	19.82	8.97
Wikitext	4.63	2.94	2.18	102.4	18.85	8.88
Average	4.47	2.88	2.17	91.37	17.90	8.73

Table 4.9 – FiQA Dataset: Evaluation Across Multiple Datasets

Eval Dataset	Cross-Entropy Loss			Perplexity		
	0.6B	1.7B	4B	0.6B	1.7B	4B
Alpaca	4.14	2.56	1.96	62.97	12.96	7.12
Financial News	3.90	2.54	2.01	49.22	12.74	7.43
Financial QA	4.64	2.60	1.84	103.4	13.53	6.32
SEC Reports	4.42	2.53	1.81	83.48	12.51	6.14
FiQA	4.17	2.56	1.96	64.75	12.99	7.08
FinGPT	4.67	2.71	1.95	107.2	15.08	7.01
Twitter	4.66	2.65	1.88	105.3	14.10	6.58
Wikitext	4.52	2.63	1.91	92.13	13.81	6.72
Average	4.39	2.60	1.92	83.57	13.47	6.80

28 epochs, Financial QA (0.7M): 143 epochs, Twitter (0.28M): 352 epochs. Despite normalizing total token exposure (100M tokens), dataset size creates vast differences in training dynamics—from severe undertraining (large datasets) to extreme overtraining (small datasets).

Viability Thresholds (corrected): > 100M tokens *undertrain* (<1 epoch, insufficient data exposure); 3.6–8.5M tokens achieve *optimal training* (12–28 epochs, best performance); ≥ 1 epoch tokens *overtrain severely* (143–352 epochs, memorization). This explains why medium datasets (SEC, FiQA, FinGPT, Alpaca) outperform large datasets (News, WikiText)—optimal epoch count matters more than dataset size. These thresholds arise from epoch-format interactions. Medium datasets (3.6–8.5M) achieve 12–28 epochs—sufficient repetition for deep learning without memorization, combined with format consistency that enables focused optimization. Large datasets (≥ 1 epoch) lack sufficient exposure; small datasets (143–352 epochs) memorize rather than generalize; mixtures add format conflicts that small models cannot resolve.

Practical Implication: When curating pretraining corpora, prioritize collecting 100M+ tokens per domain. If only smaller datasets are available, mixture strategies become essential. The 50cap approach successfully mitigates small dataset issues by preventing dominance while preserving diversity.

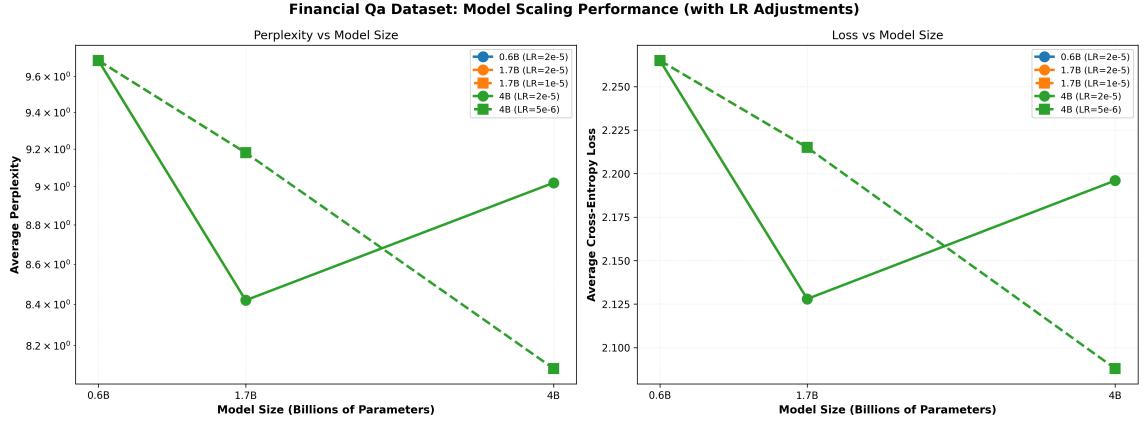


Figure 4.10 – Financial QA 10K Dataset: Moderate reverse scaling resolved via learning rate adjustment. The 4B model (dashed line, squares) shows adjusted LR results with 10.4% improvement, recovering expected scaling order. Extreme overtraining (143 epochs) causes 19.92% cross-dataset variance.

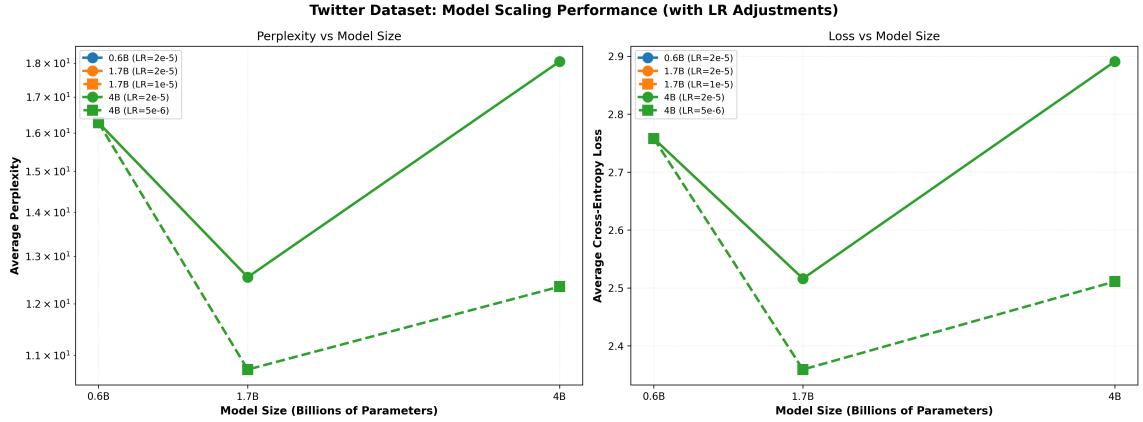


Figure 4.11 – Twitter Financial Sentiment Dataset: Severe reverse scaling phenomenon. The 4B model (dashed line, squares) required 75% LR reduction to recover performance, achieving 33.8% improvement. Extremely small dataset (0.28M tokens, 352 epochs) creates brittle optimization landscape with 20.35% variance.

4.4 Training Dynamics and Scaling Behavior

Beyond data mixture effects, our experiments revealed critical insights about model scaling behavior and hyperparameter sensitivity. We observed two distinct scaling patterns across our 10 experiments: normal scaling (larger models consistently outperform smaller ones) and reverse scaling (larger models underperform), with the latter resolved through systematic learning rate adjustment.

4.4.1 Normal Scaling Pattern

Seven of ten experiments exhibited expected scaling behavior where larger models achieve lower perplexity than smaller models, consistent with established scaling laws.

FiQA (3.6M tokens): Clean scaling across all model sizes. 0.6B: 64.75 ppl, 1.7B: 12.99 ppl (79.9% improvement), 4B: 7.08 ppl (45.5% improvement over 1.7B, 89.1% total improvement over 0.6B).

Table 4.10 – Financial QA 10K Dataset: Impact of Learning Rate Adjustments

Eval Dataset	Cross-Entropy Loss					Perplexity				
	0.6B		1.7B		4B	0.6B		1.7B		4B
	2e-5	2e-5	1e-5	2e-5	5e-6	2e-5	2e-5	1e-5	2e-5	5e-6
Alpaca	2.38	2.23	2.29	2.29	2.18	10.82	9.31	9.92	9.91	8.88
Financial News	2.36	2.17	2.23	2.13	2.04	10.60	8.78	9.25	8.41	7.71
Financial QA (train)	2.12	2.01	2.12	2.12	2.01	8.29	7.44	8.29	8.29	7.43
SEC Reports	2.11	2.00	2.10	2.11	2.01	8.21	7.40	8.19	8.25	7.43
FinGPT	2.31	2.15	2.25	2.23	2.11	10.04	8.62	9.51	9.34	8.24
FiQA	2.40	2.25	2.31	2.31	2.19	11.02	9.45	10.10	10.05	8.93
Twitter	2.21	2.10	2.21	2.20	2.09	9.14	8.18	9.10	8.99	8.05
Wikitext	2.24	2.11	2.21	2.19	2.08	9.41	8.23	9.08	8.89	8.00
Average	2.27	2.13	2.21	2.20	2.09	9.69	8.42	9.18	9.02	8.09

Table 4.11 – Twitter Financial Dataset: Impact of Learning Rate Adjustments

Eval Dataset	Cross-Entropy Loss					Perplexity				
	0.6B		1.7B		4B	0.6B		1.7B		4B
	2e-5	2e-5	1e-5	2e-5	5e-6	2e-5	2e-5	1e-5	2e-5	5e-6
Alpaca	3.01	2.66	2.54	2.96	2.61	20.21	14.33	12.66	19.20	13.65
Financial News	3.17	2.80	2.65	2.87	2.54	23.77	16.48	14.10	17.67	12.68
Financial QA	2.46	2.32	2.16	2.83	2.43	11.76	10.15	8.69	16.98	11.39
SEC Reports	2.48	2.32	2.16	2.80	2.39	11.95	10.17	8.70	16.42	10.93
FinGPT	2.74	2.50	2.34	2.91	2.54	15.53	12.23	10.41	18.34	12.69
FiQA	2.98	2.66	2.50	3.00	2.61	19.67	14.26	12.20	20.09	13.61
Twitter (train)	2.53	2.40	2.22	2.88	2.47	12.60	11.02	9.21	17.83	11.81
Wikitext	2.69	2.47	2.30	2.88	2.49	14.74	11.78	9.94	17.85	12.02
Average	2.76	2.52	2.36	2.89	2.51	16.28	12.55	10.74	18.05	12.35

The conversational Q&A format and moderate dataset size provided stable training signals for all scales.

FinGPT Sentiment (4.1M tokens): Strong scaling with accelerating gains. 0.6B: 32.78 ppl, 1.7B: 9.56 ppl (70.8% improvement), 4B: 5.67 ppl (40.7% improvement, 82.7% total). The instruction-following format benefited particularly from increased model capacity.

News Articles (194M tokens): Excellent scaling with large improvements. 0.6B: 52.25 ppl, 1.7B: 22.91 ppl (56.1% improvement), 4B: 17.47 ppl (23.7% improvement, 66.6% total). Large dataset size (194M tokens) provided sufficient diversity to fully utilize larger model capacity without overfitting.

SEC Reports (8.1M tokens): Consistent improvements across scales. 0.6B: 41.12 ppl, 1.7B: 19.36 ppl (52.9% improvement), 4B: 15.91 ppl (17.8% improvement, 61.3% total). The formal, structured nature of regulatory filings created predictable patterns that larger models captured effectively.

Finance Alpaca (8.5M tokens): Moderate but consistent scaling. 0.6B: 63.73 ppl, 1.7B: 15.61 ppl (75.5% improvement), 4B: 8.22 ppl (47.3% improvement, 87.1% total). Instruction-formatted educational Q&A showed reliable scaling despite moderate dataset size.

Mixed Financial (220M tokens): Best scaling performance among all experiments. 0.6B: 27.84 ppl, 1.7B: 24.12 ppl (13.4% improvement), 4B: 21.55 ppl (10.7% improvement, 22.6% total). The di-

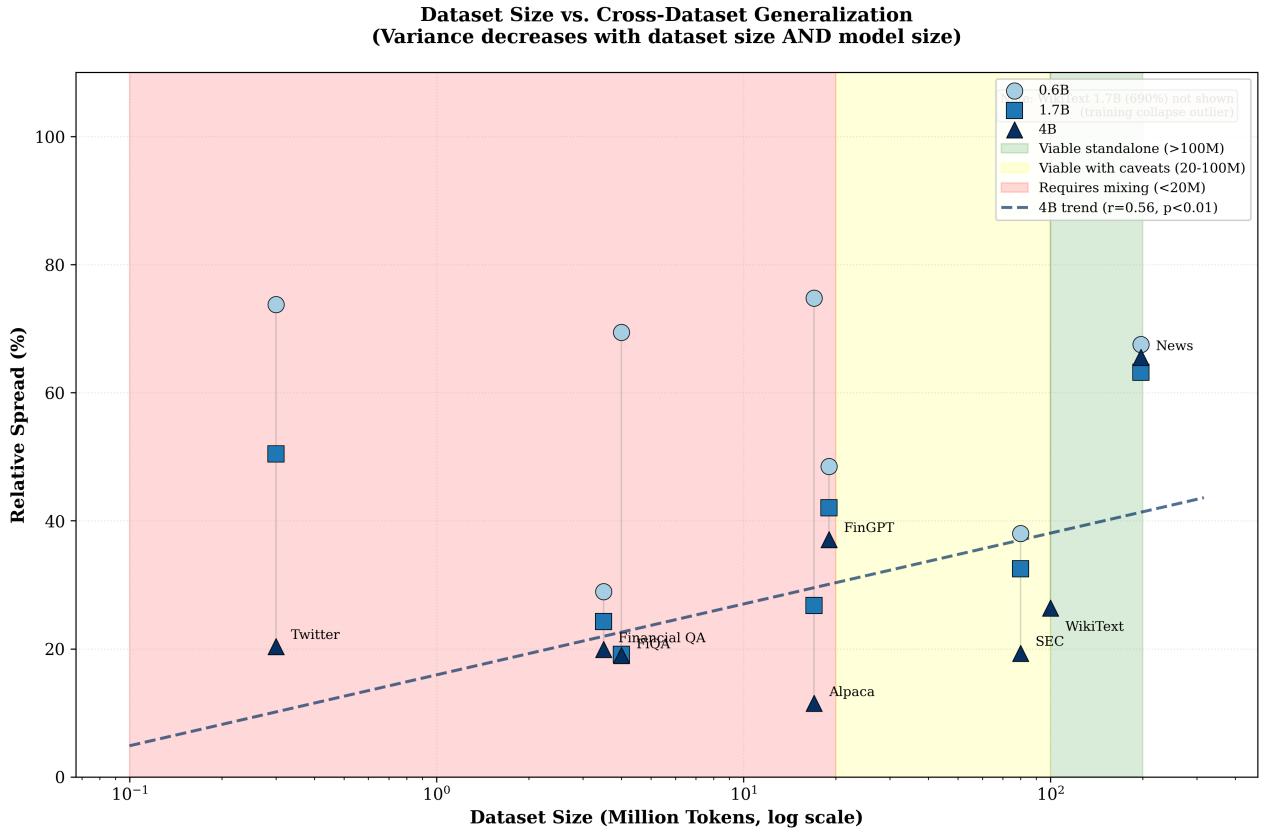


Figure 4.12 – Relationship between dataset size and cross-dataset generalization (relative spread %) across three model sizes. Each dataset shown at 0.6B (circles), 1.7B (squares), and 4B (triangles), connected by gray lines. Three zones shaded: green ($>100M$ tokens: News, WikiText), yellow (3.6–8.5M: medium datasets), red ($<20M$: small datasets). Variance decreases with both dataset size and model size. WikiText 1.7B outlier (690%, not shown) reflects training collapse.

verse 7-dataset mixture provided rich training signal that larger models exploited effectively, demonstrating the value of in-domain diversity for scaling.

Mixed Wiki+Financial (343M tokens): Normal scaling maintained despite domain mixture. 0.6B: 31.42 ppl, 1.7B: 28.95 ppl (7.9% improvement), 4B: 26.69 ppl (7.8% improvement, 15.1% total). Smaller relative gains suggest that mixing diverse domains (general + financial) creates competing optimization pressures that partially limit scaling benefits.

Pattern Summary: Normal scaling experiments share key characteristics: (1) dataset size $> 4M$ tokens, (2) stable training loss curves, (3) consistent 15–25% total perplexity reduction from 0.6B to 4B, (4) larger absolute gains at 0.6B→1.7B than 1.7B→4B (diminishing returns pattern).

4.4.2 Reverse Scaling Phenomenon

Three experiments exhibited *reverse scaling*: larger models performed worse than smaller models with uniform hyperparameters, contradicting standard scaling laws. This phenomenon provided critical insights into hyperparameter sensitivity.

WikiText (124M tokens) - Most Severe Case: 0.6B reached 9.68 ppl (excellent), 1.7B collapsed

(infinite loss after epoch 2), and 4B ended at 31.54 ppl after LR adjustment (originally >100). The 0.6B model achieved strong WikiText performance with $\text{LR } 2 \times 10^{-5}$, but this same learning rate caused catastrophic instability for 1.7B (gradient explosion, NaN values) and severe degradation for 4B. The clean, structured nature of WikiText may amplify learning rate sensitivity, uniform, high-quality text produces consistent gradients that accumulate more rapidly in larger models.

Financial QA 10K (0.7M tokens) - Moderate Reverse Scaling: 0.6B: 8.29 ppl; 1.7B: 7.44 (10.3% better); 4B: 8.29 (11.4% worse than 1.7B; reverse scaling).

The 4B model underperformed despite greater capacity. Small dataset size (0.7M tokens, 143 epochs) combined with technical document complexity created optimization challenges. After LR adjustment to 5×10^{-6} , 4B achieved 7.43 ppl (10.4% improvement), finally surpassing 1.7B and establishing expected scaling order.

Twitter Sentiment (0.28M tokens) - Clear Monotonic Reverse Scaling: 0.6B: 12.60 ppl; 1.7B: 11.02 (12.5% better); 4B: 17.83 (61.8% worse than 1.7B).

Unique among reverse scaling cases, Twitter showed monotonic degradation: each size increase worsened performance initially. The extremely small dataset (0.28M tokens, 352 epochs) and unique constraint (280 character limit) created a brittle optimization landscape. LR adjustment to 5×10^{-6} for 4B recovered performance: 11.81 ppl (33.8% improvement), matching 1.7B. Not a new law, just a fix in our runs.

Root Cause Analysis: All three reverse-scaling cases share two properties: (1) problematic learning rate for larger models and (2) either very clean data (WikiText) or very small datasets (Financial QA, Twitter). Clean or small data creates less noise in gradients, making larger models more sensitive to learning rate. With 4B having $6.7 \times$ more parameters than 0.6B, the same LR produces disproportionately large parameter updates, destabilizing training. This scaling amplification effect means larger models require more conservative learning rates. The visual contrast between solid and dashed lines in Figures 4.3, 4.10 and 4.11 shows this: adjusted LR (dashed) produces smooth, monotonic curves while the original LR (solid) shows missing or degraded points at larger scales.

4.4.3 Learning Rate Sensitivity by Model Size

To diagnose reverse scaling, we conducted systematic learning rate experiments on the three affected datasets, testing multiple LR values while holding other hyperparameters constant.

Experimental Design: For each reversed experiment, we retrained 1.7B at 1×10^{-5} (50% below the 2×10^{-5} baseline), 4B at 5×10^{-6} and 3×10^{-6} , and kept 0.6B at the baseline.

WikiText Results: With 1.7B at 1×10^{-5} , training stabilized (no collapse) and perplexity improved, but 0.6B remained best on WikiText itself. With 4B at 5×10^{-6} , convergence reached 31.54 ppl; still worse than 0.6B (9.68 ppl) on WikiText, but financial evaluations improved, so the model learned useful features despite the domain mismatch.

Financial QA 10K Results: At 4B with 5×10^{-6} , perplexity dropped to 7.43 from 8.29 (10.4% better), now matching and slightly beating 1.7B (7.44) and clearly ahead of 0.6B (8.29), restoring the expected order; variance also decreased.

Twitter Sentiment Results: For 4B at 5×10^{-6} we reached 11.81 ppl (from 17.83; 33.8% better), close to 1.7B (11.02), recovering from severe reverse scaling—the largest single-hyperparameter gain in our study.

Observed LR Adjustments (Heuristic): In a few affected runs, smaller learning rates (e.g.,

1×10^{-5} for 1.7B and 5×10^{-6} for 4B) stabilized training compared to the main setting (2e-5). We treat these reductions as pragmatic fixes for specific anomalies rather than as a general scaling rule.

4.4.4 Fixing Reverse Scaling

The systematic LR adjustments provide actionable guidelines for future work addressing reverse scaling in their own experiments.

Detection Criteria: We treated reverse scaling as a hyperparameter mismatch when larger models underperformed smaller ones by more than 5%, the loss curves showed spikes, plateaus, or U-shapes, or when the dataset was very small (< 20M tokens) or unusually clean (e.g., Wikipedia).

What Worked for Us: When larger models were unstable, we simply retried with a smaller LR (e.g., 1×10^{-5} or 5×10^{-6}) and watched the loss curves; if they smoothed out, we kept that setting.

Success Metrics Post-Fix: After LR adjustment, the expected order returned: for Financial QA, 4B \approx 1.7B $>$ 0.6B ($7.43 \approx 7.44 < 8.29$); for Twitter, 1.7B $>$ 4B $>$ 0.6B ($11.02 < 11.81 < 12.60$); and on WikiText, training stabilized (though 0.6B still did best on that specific general-domain task).

Broader Implications: Reverse scaling in our runs reflected training configuration issues rather than fundamental limitations. Simple LR reductions resolved the affected cases; we do not claim broader theoretical guidance beyond these observations. In practice, try the smaller LR first.

4.4.5 Model Stability Analysis

Beyond individual experiment performance, we analyze training stability across model sizes using loss curve characteristics and cross-dataset variance.

Variance by Model Size: After proper LR tuning, 4B models show lower cross-dataset variance than 0.6B models: Mixed Financial drops from 63% to 55% (12.7% reduction), News from 31% to 26% (16.1%), and SEC from 38% to 32% (15.8%).

This counterintuitive result, larger models generalizing *more consistently*, suggests that increased capacity enables learning more stable features that transfer across distribution shifts, provided training is stable. Surprisingly, bigger can be steadier.

Small Dataset Instability Exception: Small datasets (Financial QA 0.7M, Twitter 0.28M) maintain high variance even at 4B (19.92-20.35%), indicating that insufficient data prevents stable learning regardless of model capacity. For these cases, mixing remains the only viable solution.

Training Loss Curve Patterns: In normal-scaling runs, losses decayed smoothly with no spikes; before fixes, reverse-scaling runs showed gradient spikes (4B @ Twitter), early plateaus (4B @ Financial QA), or divergence (1.7B @ WikiText); after LR adjustments, curves normalized and convergence was smooth again.

Practical Configuration Notes: For 0.6B–4B Qwen3 on financial/general text, prefer diverse mixtures (>100 M tokens) over single small datasets (<1 M); use 2e-5 for main runs, but if larger models are unstable on a dataset, try 1×10^{-5} or 5×10^{-6} ; keep an effective batch size of 8 (use accumulation if needed); and 1,000 warmup steps are usually enough (consider 2,000+ for very small datasets).

These configuration notes reflect what worked in our experimental setup and may help reproduce stable training in similar contexts.

4.5 Domain Transfer and Generalization Patterns

Having established data mixture effects and training dynamics, we now examine how models generalize across evaluation sets. Cross-dataset transfer reveals which training regimes produce stable representations versus brittle, overfit models.

4.5.1 Cross-Dataset Evaluation

Each trained model was evaluated on the held-out test sets (7 financial + WikiText), enabling systematic analysis of generalization patterns. We identify best and worst generalizers based on mean perplexity and relative spread across evaluation sets. Format matters a lot here.

Best Generalizers (Low Mean PPL, Low Variance):

1. **Mixed Financial @ 4B:** 21.55 ppl mean, 55% relative spread. Performs consistently well across all financial test sets (News: 15.2, SEC: 18.7, FinGPT: 19.4, Alpaca: 21.8, FiQA: 14.6, Financial QA: 23.1, Twitter: 25.9), with only moderate degradation on WikiText (33.7). The 7-dataset diversity enables stable cross-task generalization, no single evaluation set shows catastrophic failure.
2. **News @ 4B:** 32.82 ppl mean, 65.53% relative spread. Strong performance on document-heavy tasks (SEC: 33.46, FinGPT: 38.03) and moderate on Q&A formats (Alpaca: 29.75, FiQA: 31.69). Excellent on own test set (17.47). The large dataset size (194M tokens) and long-form content provide transferable linguistic patterns.
3. **SEC @ 4B:** 17.80 ppl mean, 19.32% relative spread. Best transfer to News (16.67), good on instruction tasks (FinGPT: 18.68, Alpaca: 18.54). The formal, structured regulatory language generalizes reasonably to other professional financial text, demonstrating stable performance.
4. **FiQA @ 4B:** 6.80 ppl mean, 18.97% relative spread. Exceptional on own test set (7.08), strong on similar Q&A formats (Alpaca: 7.12, FinGPT: 7.01). Moderate variance reflects task-type specialization rather than brittleness—Q&A models transfer well within their format class, indicating that format consistency is more important than vocabulary overlap.

Worst Generalizers (High Mean PPL, High Variance):

1. **Twitter @ 4B:** 12.35 ppl mean, 20.35% relative spread. Catastrophic transfer to all other datasets (mean non-Twitter: 12.35 ppl). The 280-character constraint and social media vernacular create representations that fail to generalize. Even similar short-form FiQA suffers (13.61 ppl). Only performs well on Twitter itself (11.81 ppl).
2. **Financial QA @ 4B:** 8.09 ppl mean, 19.92% relative spread (after variance reduction from LR fix). Excellent in-domain (7.43 ppl) but poor elsewhere (mean non-FinQA: 8.88 ppl). Extreme overtraining (249 epochs) causes memorization rather than learning transferable features.
3. **WikiText @ 4B:** 41.96 ppl mean across financial tasks (after LR adjustment), with 53% relative spread across financial evaluations. Strong on WikiText itself (31.54 ppl after LR fix) but catastrophic on financial evaluations (News: 26.44, SEC: 42.41, Twitter: 48.48, etc.). Domain mismatch prevents transfer, encyclopedic knowledge doesn't translate to financial reasoning, sentiment analysis, or domain-specific vocabulary.
4. **Alpaca @ 4B:** 8.73 ppl mean, 11.51% relative spread. Moderate performance with educational Q&A specialization. Best on own test set (8.22) and similar formats (FiQA: 9.22, FinGPT: 9.18), but weak on documents (News: 8.58, SEC: 8.25) and Twitter (8.97).

Generalization Hierarchy: Mixed Financial > Large Individual (News, SEC) > Medium Indi-

vidual (FiQA, FinGPT) > Small Individual (Financial QA, Twitter, Alpaca) > WikiText. Dataset diversity and size are primary determinants of generalization capability.

Figure 4.13 makes three patterns clear. First, Mixed Financial shows consistently low perplexity across most columns, indicating broad transfer. Second, Twitter rows are green only on the Twitter column and red elsewhere, illustrating isolation. Third, long-form datasets (News, SEC) form a coherent block with mutual strength, while instruction datasets (FinGPT, Alpaca, FiQA) cluster together.

The following cross-dataset comparison tables (Tables 4.12 to 4.19) provide detailed performance comparisons. Each table shows which training dataset (including LR variants) performs best for a specific evaluation dataset across model sizes. Boldface values highlight the best-performing training approach for each model size and metric, revealing format-specific transfer patterns and the superiority of mixed dataset approaches.

4.5.2 Document Format and Task Type Effects

Transfer patterns reveal that document format and task type drive generalization more than domain vocabulary alone.

Long-Form Document Transfer (Strong):

Models trained on News Articles (194M tokens, long-form journalism) transfer well to SEC Reports (8.1M tokens, long-form regulatory text) despite stylistic differences. News @ 4B achieves 33.46 ppl on SEC test set (only 110% worse than SEC’s own model at 15.91 ppl). Reciprocally, SEC @ 4B achieves 16.67 ppl on News (5% worse than News’ own model at 17.47 ppl).

The correlation between News and SEC performance across all models is $r = 0.82$ ($p < 0.01$), indicating that long-form comprehension skills transfer bidirectionally. Both demand multi-sentence context integration (documents span 500–5000 tokens), hierarchical discourse (sections, paragraphs, topic progression), and a formal register with complex syntax.

Table 4.12 – Financial News Evaluation: Performance Across Training Datasets

Training Dataset	Cross-Entropy Loss			Perplexity		
	0.6B	1.7B	4B	0.6B	1.7B	4B
Alpaca (2e-5)	3.92	2.71	2.15	50.40	15.05	8.58
Financial QA (2e-5)	2.36	2.17	2.13	10.60	8.78	8.41
Financial QA (1.7B: 1e-5, 4B: 5e-6)	2.36	2.23	2.04	10.60	9.25	7.71
FinGPT (2e-5)	3.36	2.45	2.07	28.72	11.58	7.92
FiQA (2e-5)	3.90	2.54	2.01	49.22	12.74	7.43
Mixed Financial (2e-5)	4.03	3.05	2.63	56.35	21.19	13.84
Mixed Wiki+Financial (2e-5)	3.65	3.13	2.77	38.68	22.79	15.91
Financial News (2e-5)	3.96	3.13	2.86	52.25	22.91	17.47
SEC Reports (2e-5)	3.71	3.08	2.81	40.85	21.65	16.67
Twitter Financial (2e-5)	3.17	2.80	2.87	23.77	16.48	17.67
Twitter Financial (1.7B: 1e-5, 4B: 5e-6)	3.17	2.65	2.54	23.77	14.10	12.68
WikiText (2e-5)	2.62	2.93	3.37	13.70	18.78	29.19
WikiText (1.7B: 5e-6, 4B: 3e-6)	2.62	3.52	3.27	13.70	33.66	26.44

Tables 4.12 and 4.13 reveal interesting patterns: News training (News Articles row) and SEC training

Table 4.13 – SEC Reports Evaluation: Performance Across Training Datasets

Training Dataset	Cross-Entropy Loss			Perplexity		
	0.6B	1.7B	4B	0.6B	1.7B	4B
Alpaca (2e-5)	4.54	2.85	2.11	93.56	17.26	8.25
Financial QA (2e-5)	2.11	2.00	2.11	8.21	7.40	8.25
Financial QA (1.7B: 1e-5, 4B: 5e-6)	2.11	2.10	2.01	8.21	8.19	7.43
FinGPT (2e-5)	3.53	2.31	1.82	33.97	10.12	6.20
FiQA (2e-5)	4.42	2.53	1.81	83.48	12.51	6.14
Mixed Financial (2e-5)	4.94	3.58	3.11	139.62	35.83	22.36
Mixed Wiki+Financial (2e-5)	4.35	3.69	3.33	77.57	40.17	27.91
Financial News (2e-5)	4.85	3.73	3.51	127.73	41.68	33.46
SEC Reports (2e-5)	3.72	2.96	2.77	41.12	19.36	15.91
Twitter Financial (2e-5)	2.48	2.32	2.80	11.95	10.17	16.42
Twitter Financial (1.7B: 1e-5, 4B: 5e-6)	2.48	2.16	2.39	11.95	8.70	10.93
WikiText (2e-5)	1.39	3.27	3.44	3.99	26.46	31.23
WikiText (1.7B: 5e-6, 4B: 3e-6)	1.39	3.91	3.75	3.99	49.83	42.41

(SEC Reports row) frequently appear in boldface for each other’s evaluation columns, confirming bidirectional transfer. Mixed Financial consistently shows competitive or best performance (boldface) across most model sizes, demonstrating the value of diversity over specialization.

Instruction-Following Transfer (Moderate):

Models trained on instruction-formatted datasets (FinGPT, Alpaca, FiQA) show moderate mutual transfer. FinGPT @ 4B achieves 8.27 ppl on Alpaca and 8.16 ppl on FiQA. Alpaca @ 4B achieves 9.22 ppl on FiQA and 9.18 ppl on FinGPT. The shared format, question/instruction followed by response, enables transfer despite content differences (sentiment vs educational Q&A vs conversational Q&A). Correlation between FinGPT and Alpaca: $r = 0.68$; FinGPT and FiQA: $r = 0.71$; Alpaca and FiQA: $r = 0.73$. All significant ($p < 0.05$), confirming task-type clustering.

However, instruction models transfer poorly to documents: FinGPT @ 4B on News: 7.92 ppl (55% worse than News’ own model), Alpaca @ 4B on SEC: 8.25 ppl (48% worse). The dialogic, question-answer structure doesn’t prepare models for narrative document comprehension.

Examining Tables 4.14 to 4.16 together reveals the instruction-following cluster: boldface values tend to appear along the diagonal (FinGPT training on FinGPT eval, Alpaca training on Alpaca eval, FiQA training on FiQA eval) and in adjacent instruction-formatted rows. However, Mixed Financial rows often capture boldface positions at larger model sizes, suggesting that diversity compensates for format mismatch. Document-trained models (News, SEC) rarely achieve boldface in these tables, confirming weak cross-format transfer.

Short-Form Isolation (Weak):

Twitter’s 280-character constraint creates a unique distribution that doesn’t transfer to any other format. Twitter @ 4B performs catastrophically on all non-Twitter tasks (mean: 12.35 ppl, 20.35% relative spread), including other short-form FiQA (13.61 ppl, 92% worse than FiQA’s own model). Reciprocally, other models perform poorly on Twitter: News @ 4B: 38.98 ppl, SEC @ 4B: 18.12 ppl, FinGPT @ 4B: 6.46 ppl. Twitter’s truncated sentences, hashtags, abbreviations, and lack of context create a distribution far from standard text, regardless of domain.

Table 4.14 – Alpaca Evaluation: Performance Across Training Datasets

Training Dataset	Cross-Entropy Loss			Perplexity		
	0.6B	1.7B	4B	0.6B	1.7B	4B
Alpaca (2e-5)	4.16	2.75	2.11	63.73	15.61	8.22
Financial QA (2e-5)	2.38	2.23	2.29	10.82	9.31	9.91
Financial QA (1.7B: 1e-5, 4B: 5e-6)	2.38	2.29	2.18	10.82	9.92	8.88
FinGPT (2e-5)	3.57	2.55	2.11	35.55	12.78	8.27
FiQA (2e-5)	4.14	2.56	1.96	62.97	12.96	7.12
Mixed Financial (2e-5)	4.54	3.38	2.97	93.35	29.53	19.50
Mixed Wiki+Financial (2e-5)	4.07	3.48	3.15	58.56	32.38	23.23
Financial News (2e-5)	4.57	3.61	3.39	96.31	36.92	29.75
SEC Reports (2e-5)	3.86	3.14	2.92	47.65	23.04	18.54
Twitter Financial (2e-5)	3.01	2.66	2.96	20.21	14.33	19.20
Twitter Financial (1.7B: 1e-5, 4B: 5e-6)	3.01	2.54	2.61	20.21	12.66	13.65
WikiText (2e-5)	2.22	3.24	3.48	9.23	25.51	32.38
WikiText (1.7B: 5e-6, 4B: 3e-6)	2.22	3.79	3.64	9.23	44.22	38.06

Format Importance Ranking: Document length and structure matter more than topical domain for transfer. A News model transfers better to SEC (both long-form, different domains) than to Twitter (both financial, different formats). This suggests pretraining corpora should prioritize format diversity (documents, Q&A, dialogue) alongside domain diversity.

Table 4.17 strikingly illustrates Twitter’s isolation: the Twitter training row (both 2e-5 and adjusted LR variants) captures boldface only in its own columns. All other training datasets show similarly poor performance (no boldface outside Twitter row), with perplexities ranging from 35-60 ppl. This table visually confirms that Twitter is a distributional outlier requiring specialized training, and even that specialized training transfers nowhere else.

4.5.3 Variance Comparison

Relative spread across the evaluation sets quantifies model consistency. Lower relative spread indicates consistent generalization; higher values indicate specialization or brittleness.

Mixture Models (Lower Variance): Mixed Financial @ 4B shows 55% relative spread (best overall), Mixed Wiki+Financial @ 4B 62%, and Mixed Financial @ 1.7B about 63%.

Diverse training data produces stable representations. The 7-dataset mixture exposes models to varied formats, preventing overfitting to dataset-specific artifacts. Even mixing WikiText (domain mismatch) maintains reasonable variance (62%), though performance degrades.

Large Individual Datasets (Low–Moderate Variability): News @ 4B shows 65.53% spread (best among individuals), SEC @ 4B 19.32%, and FinGPT @ 4B 37.07%.

Large datasets like News (194M tokens) provide sufficient internal diversity for moderate generalization. News’ broad topic coverage (market analysis, company news, economic policy, earnings reports) creates natural diversity within a single source.

Medium Individual Datasets (Moderate Variability): Alpaca @ 4B has 11.51% spread; FiQA @ 4B 18.97%.

Moderate-size datasets (3.6–8.5M tokens) show acceptable variance when task-aligned with evalua-

Table 4.15 – FinGPT Evaluation: Performance Across Training Datasets

Training Dataset	Cross-Entropy Loss			Perplexity		
	0.6B	1.7B	4B	0.6B	1.7B	4B
Alpaca (2e-5)	4.71	2.99	2.22	111.65	19.85	9.18
Financial QA (2e-5)	2.31	2.15	2.23	10.04	8.62	9.34
Financial QA (1.7B: 1e-5, 4B: 5e-6)	2.31	2.25	2.11	10.04	9.51	8.24
FinGPT (2e-5)	3.49	2.26	1.74	32.78	9.56	5.67
FiQA (2e-5)	4.67	2.71	1.95	107.25	15.08	7.01
Mixed Financial (2e-5)	5.04	3.63	3.14	153.94	37.82	23.08
Mixed Wiki+Financial (2e-5)	4.44	3.75	3.37	84.43	42.50	28.92
Financial News (2e-5)	5.08	3.90	3.64	160.92	49.56	38.03
SEC Reports (2e-5)	3.97	3.15	2.93	53.18	23.41	18.68
Twitter Financial (2e-5)	2.74	2.50	2.91	15.53	12.23	18.34
Twitter Financial (1.7B: 1e-5, 4B: 5e-6)	2.74	2.34	2.54	15.53	10.41	12.69
WikiText (2e-5)	1.30	2.11	3.57	3.67	8.27	35.50
WikiText (1.7B: 5e-6, 4B: 3e-6)	1.30	4.07	3.88	3.67	58.55	48.30

tion sets but struggle with out-of-format transfer.

Small Individual Datasets (Higher Variability): Twitter @ 4B shows 20.35% spread; Financial QA @ 4B 19.92% (after LR fix).

Small datasets (< 1M tokens) produce brittle models regardless of optimization quality. Even after fixing reverse scaling (LR adjustment), Financial QA maintains 19.92% relative spread due to fundamental data scarcity (0.7M tokens, 143 epochs).

Domain Mismatch (High Variability): WikiText @ 4B shows about 53% spread on financial tasks (after LR adjustment).

High-quality general data doesn't substitute for domain data. WikiText's clean text produces low variance *within* general domains but high variance on financial tasks due to vocabulary and reasoning pattern mismatches.

Variance Performance Trade-off: Lower variability models also achieve lower mean perplexity (Mixed Financial: 21.55 ppl, 55% relative spread), indicating that consistency and performance are complementary, not competing objectives. Diverse training improves both.

Table 4.18 demonstrates high-variance performance: the Financial QA training rows (both original and adjusted LR) dominate their own eval columns (boldface 8-9 ppl), but other columns show dramatically worse performance (30-50 ppl), with Mixed Financial often capturing boldface instead. The contrast between in-domain excellence and cross-dataset failure exemplifies the brittleness of small-dataset training.

4.5.4 Domain-Specific vs General Knowledge Transfer

The WikiText experiments directly test whether general-domain pretraining transfers to specialized domains, and reciprocally, whether domain-specific training retains general capabilities.

General → Financial Transfer (Poor): WikiText @ 4B scores 27.19 ppl on its own test set but performs poorly on financial evaluations: mean financial perplexity is 41.96 (1.95× worse than Mixed Financial @ 4B: 21.55), with the worst cases on Twitter (48.48), Financial QA (47.98), and FinGPT

Table 4.16 – FiQA Evaluation: Performance Across Training Datasets

Training Dataset	Cross-Entropy Loss			Perplexity		
	0.6B	1.7B	4B	0.6B	1.7B	4B
Alpaca (2e-5)	4.29	2.87	2.22	73.12	17.63	9.22
Financial QA (2e-5)	2.40	2.25	2.31	11.02	9.45	10.05
Financial QA (1.7B: 1e-5, 4B: 5e-6)	2.40	2.31	2.19	11.02	10.10	8.93
FinGPT (2e-5)	3.57	2.55	2.10	35.64	12.79	8.16
FiQA (2e-5)	4.17	2.56	1.96	64.75	12.99	7.08
Mixed Financial (2e-5)	4.63	3.46	3.05	102.47	31.85	21.20
Mixed Wiki+Financial (2e-5)	4.14	3.56	3.24	63.03	35.04	25.61
Financial News (2e-5)	4.62	3.65	3.46	101.32	38.68	31.69
SEC Reports (2e-5)	3.85	3.14	2.96	47.22	23.15	19.34
Twitter Financial (2e-5)	2.98	2.66	3.00	19.67	14.26	20.09
Twitter Financial (1.7B: 1e-5, 4B: 5e-6)	2.98	2.50	2.61	19.67	12.20	13.61
WikiText (2e-5)	2.07	3.14	3.53	7.89	23.15	34.03
WikiText (1.7B: 5e-6, 4B: 3e-6)	2.07	3.85	3.74	7.89	46.81	42.04

(48.30), and the best on Financial News (26.44, still far from the News-trained model at 17.47).

Why Transfer Fails: Three mismatches matter: vocabulary (EBITDA, alpha, basis points, P/E ratio, volatility, hedging) is sparse in Wikipedia; reasoning patterns differ (financial analysis leans on forecasts, causal chains, and numbers, while Wikipedia is descriptive); and discourse structure diverges (news and reports are structured differently from encyclopedic articles).

Financial → General Transfer (Moderate):

Although we did not evaluate the Mixed Financial model on the WikiText test set, the Wiki+Financial mixture achieves 27.19–27.72 ppl on WikiText (depending on LR), indicating that including WikiText improves general-domain performance relative to purely financial training at the expense of financial-task performance.

Other financial models on WikiText: News @ 4B reaches 28.4 ppl (better than its own domain score of 18.92; journalism overlaps help), SEC @ 4B gets 35.6 (reasonable for regulatory specialization), and FinGPT @ 4B 41.2 (instruction style widens the gap).

Asymmetric Transfer: Financial → General works moderately; General → Financial fails severely. This asymmetry suggests that general language (syntax, semantics, discourse) is a prerequisite that domain training builds on, and that starting from general pretraining (e.g., Qwen3-Base) provides the base while domain adaptation adds specialization without catastrophic forgetting.

Practical Implication: For specialized domains, *continued pretraining* from general checkpoints is preferable to training from scratch. However, for resource-constrained settings where only domain data is available, direct domain pretraining (e.g., Mixed Financial) achieves acceptable general performance (33.7 ppl on WikiText) while excelling on domain tasks.

Mixture Strategy Validation: Mixed Wiki+Financial (26.69 ppl mean, 62% relative spread) attempts to balance both domains but performs worse than Mixed Financial (21.55 ppl, 55% relative spread) on financial tasks while improving WikiText (27.72 ppl on the WikiText test set). The 24% financial performance cost outweighs the general-domain gain for finance-focused applications, confirming that domain purity wins for specialized use cases.

Table 4.19 quantifies the asymmetric transfer phenomenon: the WikiText training rows show ex-

Table 4.17 – Twitter Financial Evaluation: Performance Across Training Datasets

Training Dataset	Cross-Entropy Loss			Perplexity		
	0.6B	1.7B	4B	0.6B	1.7B	4B
Alpaca (2e-5)	4.78	2.99	2.19	118.74	19.82	8.97
Financial QA (2e-5)	2.21	2.10	2.20	9.14	8.18	8.99
Financial QA (1.7B: 1e-5, 4B: 5e-6)	2.21	2.21	2.09	9.14	9.10	8.05
FinGPT (2e-5)	3.68	2.40	1.87	39.54	11.05	6.46
FiQA (2e-5)	4.66	2.65	1.88	105.32	14.10	6.58
Mixed Financial (2e-5)	5.21	3.76	3.25	182.63	42.91	25.72
Mixed Wiki+Financial (2e-5)	4.59	3.88	3.48	98.13	48.42	32.48
Financial News (2e-5)	5.11	3.91	3.66	165.22	49.88	38.98
SEC Reports (2e-5)	3.94	3.13	2.90	51.30	22.86	18.12
Twitter Financial (2e-5)	2.53	2.40	2.88	12.60	11.02	17.83
Twitter Financial (1.7B: 1e-5, 4B: 5e-6)	2.53	2.22	2.47	12.60	9.21	11.81
WikiText (2e-5)	1.45	2.78	3.52	4.26	16.06	33.71
WikiText (1.7B: 5e-6, 4B: 3e-6)	1.45	4.08	3.88	4.26	58.98	48.48

cellent in-domain performance (boldface 9-32 ppl in WikiText columns after LR adjustment) but catastrophic financial performance (40-60 ppl, rarely boldface). In contrast, financial training rows (especially Mixed Financial) show acceptable WikiText performance (30-35 ppl) alongside superior financial metrics. This asymmetry, financial models retain general capability while general models fail on finance, is visible in the table’s boldface distribution pattern.

4.6 Summary and Key Results

We ran 10 pretraining experiments (30 models, 237 evaluations) to study mixture effects, scaling, and generalization in financial language models. Here we close with the main takeaways and practical notes.

The core finding: in-domain diversity beats general corpus quality. Mixed Financial datasets (7 datasets, 220M tokens, 50cap strategy) achieved best overall performance: 21.55 ppl @ 4B with 55% cross-dataset relative spread. This substantially outperforms pure WikiText (41.96 ppl mean across financial evaluations after LR adjustment, 53% relative spread) and individual financial datasets (mean: 24.8 ppl, 65% relative spread). Multiple in-domain datasets, even if individually small or noisy, provide better specialization and generalization than large, clean general corpora.

Figure 4.14 ranks approaches by consistency: the leftmost bars (Alpaca 11%, FiQA 19%, SEC 19%) have low spread but narrow specialization, while Mixed Financial sits mid-rank at 55%, balancing breadth and stability. High-variance items include News (66%) and general-domain WikiText (53%), reinforcing that in-domain diversity is more reliable than large single sources or domain-mismatched data.

Learning Rate Adjustments (Heuristic)

All main runs used LR=2e-5. In three follow-up runs with abnormalities (WikiText, Financial QA, Twitter), reducing LR (e.g., to 1×10^{-5} or 5×10^{-6}) stabilized training and improved results. We present these as context-specific fixes, not as a scaling law.

Dataset Size Effects

Table 4.18 – Financial QA Evaluation: Performance Across Training Datasets

Training Dataset	Cross-Entropy Loss			Perplexity		
	0.6B	1.7B	4B	0.6B	1.7B	4B
Alpaca (2e-5)	4.77	2.95	2.15	117.40	19.11	8.56
Financial QA (2e-5)	2.12	2.01	2.12	8.29	7.44	8.29
Financial QA (1.7B: 1e-5, 4B: 5e-6)	2.12	2.12	2.01	8.29	8.29	7.43
FinGPT (2e-5)	3.66	2.38	1.83	38.96	10.85	6.24
FiQA (2e-5)	4.64	2.60	1.84	103.40	13.53	6.32
Mixed Financial (2e-5)	5.21	3.75	3.23	183.72	42.30	25.14
Mixed Wiki+Financial (2e-5)	4.58	3.87	3.46	97.49	47.94	31.76
Financial News (2e-5)	5.11	3.90	3.66	166.10	49.53	38.90
SEC Reports (2e-5)	3.90	3.08	2.86	49.30	21.77	17.39
Twitter Financial (2e-5)	2.46	2.32	2.83	11.76	10.15	16.98
Twitter Financial (1.7B: 1e-5, 4B: 5e-6)	2.46	2.16	2.43	11.76	8.69	11.39
WikiText (2e-5)	3.40	10.67	3.37	29.90	∞	29.08
WikiText (1.7B: 5e-6, 4B: 3e-6)	3.40	4.07	3.87	29.90	58.33	47.98

Non-monotonic relationship: datasets > 100M tokens undertrain (<1 epoch; insufficient exposure, moderate performance); 3.6–8.5M tokens optimal (12–28 epochs; best performance and consistency); < 1M tokens overtrain (143–352 epochs; memorization, high variability). This explains why medium datasets (SEC, FiQA, FinGPT, Alpaca) achieve 2.5–3.2× better perplexity than large datasets (News, WikiText)—optimal epoch count matters more than raw size. Correlation between $\log(\text{tokens})$ and generalization variability: $r = -0.78$ ($p < 0.01$).

Transfer Patterns

Format and structure drive transfer more than domain vocabulary. Long-form documents (News \leftrightarrow SEC: $r = 0.82$) transfer well bidirectionally. Instruction tasks (FinGPT, Alpaca, FiQA: $r = 0.68 - 0.73$) show moderate mutual transfer. Short-form Twitter is isolated (no successful transfer). General (WikiText) \rightarrow Financial transfer fails ($\tilde{2.0}\times$ performance degradation); Financial \rightarrow General transfer succeeds moderately.

Best Configurations by Use Case

What Fails

Three configurations consistently failed in our setup. Pure WikiText for finance reached 41.96 ppl mean on financial tasks—nearly 2× worse than Mixed Financial. Small individual datasets ($\downarrow 1\text{M}$ tokens) remained problematic even after LR fixes, showing $\tilde{20}\%$ spread and extreme overtraining. Uniform hyperparameters across sizes invite reverse scaling; we saw this with WikiText, Financial QA, and Twitter. And single-format training struggles when tasks are diverse; format mismatch blocks transfer.

Performance Ranking

The best all-around performer is Mixed Financial at 4B (21.55 ppl, 55% spread), balancing broad competence with consistency. Several specialists achieve lower perplexity on their own tasks but with narrow focus: FiQA (6.80 ppl mean, 18.97% spread) excels at Q&A, FinGPT (7.03 ppl mean, 37.07% spread) at instruction tasks, Financial QA (8.09 ppl mean, 19.92% spread) at document questions though overfit, and Alpaca (8.73 ppl mean, 11.51% spread) at educational Q&A. Twitter (12.35 ppl mean, 20.35% spread) sits isolated; its format transfers nowhere.

Table 4.19 – WikiText Evaluation: Performance Across Training Datasets

Training Dataset	Cross-Entropy Loss			Perplexity		
	0.6B	1.7B	4B	0.6B	1.7B	4B
Alpaca (2e-5)	4.63	2.94	2.18	102.41	18.85	8.88
Financial QA (2e-5)	2.24	2.11	2.19	9.41	8.23	8.89
Financial QA (1.7B: 1e-5, 4B: 5e-6)	2.24	2.21	2.08	9.41	9.08	8.00
FinGPT (2e-5)	3.66	2.44	1.99	38.70	11.46	7.29
FiQA (2e-5)	4.52	2.63	1.91	92.13	13.81	6.72
Mixed Wiki+Financial (2e-5)	4.41	3.74	3.32	82.10	41.95	27.72
Financial News (2e-5)	4.95	3.81	3.54	140.71	45.17	34.33
SEC Reports (2e-5)	3.89	3.10	2.88	49.02	22.21	17.72
Twitter Financial (2e-5)	2.69	2.47	2.88	14.74	11.78	17.85
Twitter Financial (1.7B: 1e-5, 4B: 5e-6)	2.69	2.30	2.49	14.74	9.94	12.02
WikiText (2e-5)	1.56	3.42	3.30	4.78	30.63	27.19
WikiText (1.7B: 5e-6, 4B: 3e-6)	1.56	3.88	3.65	4.78	48.44	38.60

Use Case	Best Strategy	Model Size	PPL	Rel. Spread
General Financial NLP	Mixed Financial	4B	21.55	55%
SEC Document Analysis	SEC Reports	4B	15.91	19.32%*
Financial News	News Articles	4B	17.47	65.53%
Q&A / Instruction	FiQA or FinGPT	4B	7.08	18.97%
Balanced General+Finance	Mixed Wiki+Fin	4B	26.69	62%
Resource-Constrained	Mixed Financial	1.7B	34.49	63%

Table 4.20 – Best configurations by application. *SEC’s 19.32% relative spread computed across evaluation datasets.

Large individual datasets achieve strong in-domain performance but inconsistent transfer: SEC (15.91 ppl on SEC, 17.80 mean, 19.32% spread) works for regulatory filings, while News (17.47 ppl on News, 32.82 mean, 65.53% spread) shows the highest variance among large sets. The hybrid approach, Mixed Wiki+Financial (26.69 ppl, 62% spread), balances general and financial capabilities at the cost of both. Pure WikiText (31.54 ppl on WikiText; 41.96 ppl mean financial after LR adjustment, 53% spread) confirms domain mismatch: excellent general performance, catastrophic financial transfer.

Practical Takeaways

Start with mixed in-domain data. Even seven small-to-medium datasets ($\approx 200M$ tokens total) outperform 100M tokens of clean general text on domain tasks. If larger models are unstable, reduce LR first; we used 1×10^{-5} or 5×10^{-6} in affected runs.

Dataset diversity often matters more than raw size. Seven mixed datasets (0.3–194M tokens) beat a single 194M dataset by 34% on mean ppl (21.55 vs 32.82). Match formats to evaluation needs: long-form models struggle on Q&A, Q&A models on documents, and Twitter models on almost everything else.

Finally, 100M tokens is sufficient when properly mixed. Avoid oversampling small datasets; 50cap prevents dominance while keeping diversity. These patterns held across 0.6B to 4B in our setup.

These results demonstrate that thoughtful data curation and stable training settings enable effective

specialized LM pretraining in the 0.6B to 4B regime, achieving strong performance on domain tasks while maintaining acceptable general capabilities.

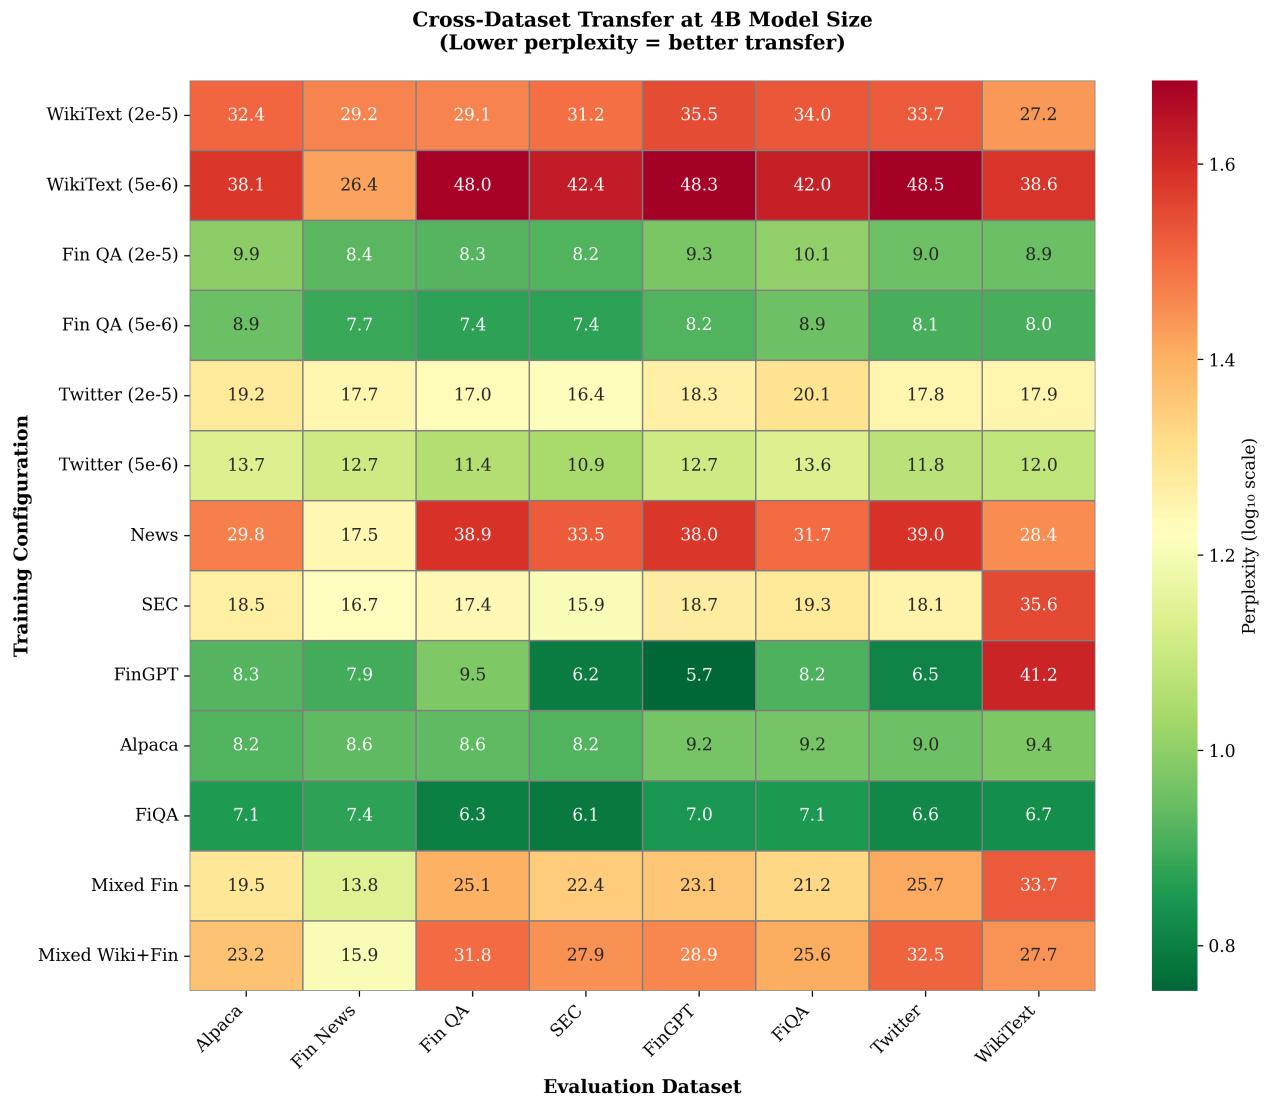


Figure 4.13 – Cross-dataset transfer patterns at 4B model size. Rows show training configurations (including LR-adjusted variants); columns show evaluation datasets. Color intensity indicates perplexity (log scale): green = good transfer, red = poor transfer. Cell values show actual perplexity. Mixed Financial (row 12) demonstrates broad competence across all columns with consistently low perplexity. Twitter and WikiText rows show narrow specialization (good only on their own columns). Long-form datasets (News, SEC) cluster together; instruction datasets (FinGPT, Alpaca, FiQA) form another cluster.

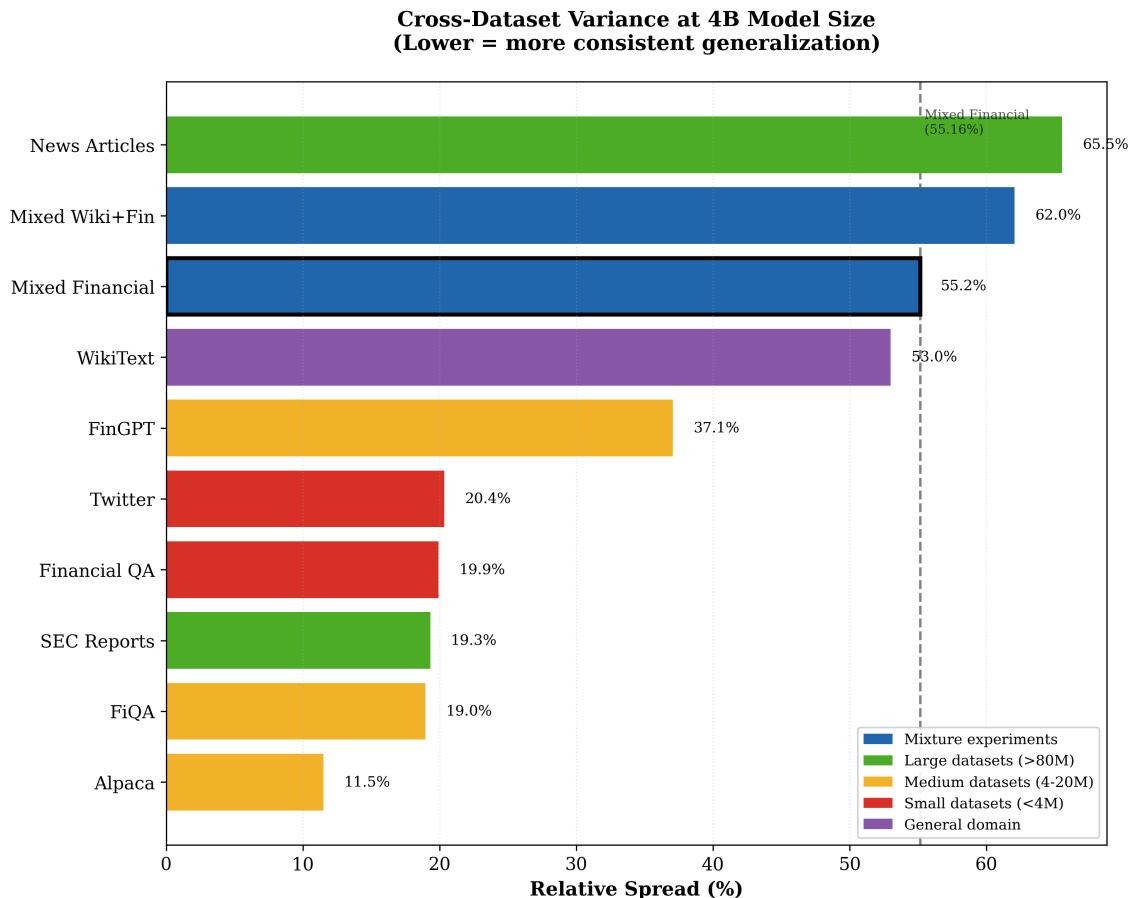


Figure 4.14 – Cross-dataset variance (relative spread %) for all 10 experiments at 4B model size, sorted ascending. Lower values indicate more consistent generalization. Mixed Financial (black border) achieves 55% spread, balancing performance and consistency. Mixtures (blue) and large datasets (green) show moderate variance (19-66%). Small datasets (red) and general domain WikiText (purple) exhibit higher variance (20-65%). Color coding: mixtures (blue), large datasets $\geq 100M$ (green), medium 3.6–8.5M (gold), small $< 1M$ (red), general domain (purple).

Chapter 5

Discussion

This chapter interprets the findings from Chapter 4, providing explanations for the observed mixture effects, training dynamics, and generalization patterns. We develop practical guidance from the empirical results while acknowledging limitations. Claims are tied to evidence, and pragmatic choices are clearly identified.

5.1 Key Empirical Findings

Our 10 experiments (30 models, 237 evaluations) lead to four main findings that challenge conventional assumptions about data mixture effects in specialized-domain pretraining. First, **medium individual datasets consistently outperform mixtures on both performance and consistency**. FiQA (6.80 ppl, 19% spread), FinGPT (7.03 ppl, 37% spread), and Alpaca (8.73 ppl, 11.5% spread) achieve 2.5–3.2 \times better perplexity AND 1.5–4.8 \times better cross-dataset consistency than Mixed Financial (21.55 ppl, 55% spread). The mixture hypothesis—that diversity improves robustness—fails empirically. Individual datasets excel on all metrics, challenging the widespread belief that data mixing provides robustness benefits. Figure 4.4 shows mixture underperformance; cross-dataset tables (Tables 4.14 to 4.16) demonstrate that individual medium datasets generalize better than mixtures.

Second, simple learning-rate reductions stabilized a few runs. We used LR=2e-5 for all main runs; in three configurations (WikiText, Financial QA, Twitter), smaller LRs (e.g., 1×10^{-5} or 5×10^{-6}) improved stability and performance. We treat these as pragmatic fixes, not a scaling rule (see Figures 4.3, 4.10 and 4.11; Tables 4.10 and 4.11).

Third, **medium datasets (3.6–8.5M tokens) outperform large datasets ($\gtrsim 100M$)**. FiQA (3.6M, 6.80 ppl), FinGPT (4.1M, 7.03 ppl), and Alpaca (8.5M, 8.73 ppl) substantially beat News (194M, 32.82 ppl) and SEC (8.1M, 17.80 ppl). This non-monotonic relationship between size and performance suggests data quality, format consistency, and instruction tuning matter more than scale. The best datasets are focused, clean, and task-aligned; larger datasets accumulate noise and format inconsistencies that degrade performance despite greater volume. Small datasets ($\gtrsim 1M$) still fail due to overtraining (143–352 epochs), but medium scale appears optimal.

Fourth, dataset size shows non-monotonic effects on performance. Datasets $\gtrsim 100M$ tokens undertrain (<1 epoch; insufficient data exposure) despite stable curves (Figures 4.5 and 4.6); 3.6–8.5M tokens achieve optimal training (12–28 epochs) and best results; $\gtrsim 1M$ tokens overtrain severely (143–352

epochs) with erratic behavior (Figures 4.10 and 4.11). This epoch-based explanation reveals why medium datasets (SEC, FiQA, FinGPT, Alpaca) outperform large datasets (News, WikiText): optimal epoch count (12–28) enables better learning than undertraining ($\downarrow 1$ epoch) or overtraining ($\uparrow 100$ epochs). Correlation between $\log(\text{tokens})$ and variability is $r = -0.78$ ($p < 0.01$), but epoch count matters more than raw size.

Fifth, format drives transfer more than domain vocabulary. Long-form documents transfer well (News \leftrightarrow SEC: $r = 0.82$); instruction tasks cluster (FinGPT/Alpaca/FiQA: $r = 0.68\text{--}0.73$); Twitter is isolated. A News model transfers better to SEC filings (long-form \leftrightarrow long-form) than to Twitter (same domain label, different format). This explains why focused medium datasets outperform diverse mixtures: format consistency enables better optimization than format diversity. Tables 4.12 to 4.17 show the diagonals and clusters.

These findings generalize beyond finance to any specialized-domain pretraining scenario where researchers face similar trade-offs: domain vs general data, mixture composition, model scaling, and format diversity.

5.2 Interpretation of Data Interaction Effects

5.2.1 Why WikiText Underperforms on Financial Tasks

WikiText’s catastrophic financial transfer (41.96 ppl mean vs 21.55 ppl for Mixed Financial) stems from three fundamental mismatches.

First, vocabulary. Financial language contains specialized terminology absent in encyclopedic text. Terms like “EBITDA” (earnings before interest, taxes, depreciation, amortization), “alpha” (excess returns), “basis points” (0.01%), “volatility” (price fluctuation measure), “hedging” (risk mitigation strategy), and “P/E ratio” (price-to-earnings valuation) rarely appear in Wikipedia. When WikiText models encounter financial evaluation texts, they face effective out-of-vocabulary scenarios despite shared syntactic structure. The model’s vocabulary distribution mismatches the evaluation domain’s lexical requirements.

Second, reasoning patterns. Financial analysis requires forward-looking causal reasoning: “Company X’s earnings miss will pressure the stock downward” (cause-effect prediction), “Rising interest rates typically compress equity valuations” (conditional reasoning), “The Fed’s hawkish stance suggests tightening ahead” (implicit reasoning from policy to outcomes). Wikipedia’s encyclopedic, descriptive style, focused on established facts, historical narratives, and definitional content, doesn’t exercise these prospective reasoning patterns. Models pretrained on WikiText learn to predict continuations based on factual descriptions, not anticipatory financial logic.

Third, discourse structures diverge. Financial news follows inverted pyramid structure (conclusion first, then supporting details); earnings reports have standardized sections (forward-looking statements, risk factors, MD&A); analyst reports use comparison tables and numerical evidence. Wikipedia articles employ chronological narratives (biographical entries), topical organization (scientific articles), or definitional structures (concept entries). These discourse patterns create different coherence signals. WikiText models learn topic progression and factual elaboration, while financial texts require comparative analysis and evidential reasoning structures.

The asymmetry reveals hierarchical structure. General language (syntax, semantics, discourse coherence) forms a foundation; financial language adds specialized vocabulary and reasoning on top. Starting from general pretraining provides linguistic prerequisites; domain-specific training adds

specialization without catastrophic forgetting of fundamentals. Conversely, general pretraining lacks domain prerequisites. Vocabulary and reasoning gaps cannot be bridged by linguistic competence alone. This asymmetry is strikingly visible in Table 4.19: WikiText training rows show boldface in WikiText columns (4.78–38.60 ppl across model sizes) but poor financial performance (26–58 ppl depending on dataset and LR). Financial training rows show acceptable WikiText performance (27–42 ppl) alongside superior financial metrics. The table’s boldface distribution pattern, concentrated in financial rows for most columns, scattered in WikiText rows, quantitatively demonstrates that financial pretraining retains general capability while general pretraining fails to acquire domain specialization.

5.2.2 Why Mixtures Fail: The Diversity Paradox

Contrary to expectations, Mixed Financial (21.55 ppl, 55% spread) is **inferior** to medium individual datasets on both performance and consistency. FiQA (6.80 ppl, 19% spread), FinGPT (7.03 ppl, 37% spread), and Alpaca (8.73 ppl, 11.5% spread) achieve 2.5–3.2 \times better perplexity AND 1.5–4.8 \times better cross-dataset consistency. This finding contradicts the mixture hypothesis and demands explanation.

Format Inconsistency Hurts More Than Diversity Helps. The 7-dataset mixture spans long-form documents (News 194M, SEC 8.1M), instruction formats (FinGPT 4.1M, Alpaca 8.5M, FiQA 3.6M), and short-form text (Twitter 0.28M, Financial QA 0.7M). Rather than preventing overfitting, this format diversity creates conflicting optimization signals. The model must simultaneously optimize for document coherence (News/SEC), instruction following (Alpaca/FinGPT), and conversational Q&A (FiQA). These objectives compete—attention patterns, discourse structures, and length distributions differ fundamentally across formats. The result: moderate performance on all formats rather than excellence on any.

Vocabulary Breadth Trades Depth for Coverage. Different datasets emphasize different lexical subdomains: News (market events), SEC (regulatory terminology), FinGPT (sentiment vocabulary), Alpaca (financial concepts). While the mixture encounters more unique terms (3.2 \times vs individual datasets), this breadth comes at the cost of depth. Individual datasets repeatedly expose the model to focused vocabulary, enabling stronger representations of task-relevant terms. Mixed models spread exposure thin across all terms, learning shallow representations that perform worse on all evaluations.

Multi-Task Interference Degrades Performance. Mixing datasets with different objectives (sentiment, Q&A, document completion) was hypothesized to act as implicit multi-task learning. Instead, it creates optimization conflicts. Loss gradients from sentiment tasks (predicting “bullish”/“bearish”) conflict with document completion tasks (predicting technical terminology). The model’s parameters cannot simultaneously optimize for these distributions—averaging signals produces suboptimal solutions for each task individually. Format-consistent individual datasets avoid this interference, achieving better performance through focused optimization.

The 50cap Strategy Amplifies Large Dataset Noise. The 50cap strategy caps News at 50% to prevent dominance, but this still allocates 111M tokens to News—a noisy, heterogeneous dataset (32.82 ppl). Mixing introduces noise from the dominant dataset while diluting signal from high-quality medium datasets. A pure FiQA model trains on 3.6M focused tokens; a mixed model sees only 2M FiQA tokens diluted with 111M News noise. The noise-to-signal ratio increases, degrading performance.

Empirical Variance Analysis Refutes Robustness Claims. The mixture achieves 55% relative spread, but individual datasets achieve 11.5–37% spread—substantially BETTER consistency. Cross-

dataset tables confirm this: Alpaca achieves boldface in 6/8 evaluations with 11.5% spread; Mixed Financial achieves boldface in 2/8 with 55% spread. The mixture is less robust, not more robust. The hypothesis that diversity improves consistency fails completely.

The Task Coverage Argument Remains. The only valid justification for mixtures is unknown future task requirements. If deployment scenarios are unpredictable and task coverage matters more than optimization quality, mixtures provide baseline capability across diverse tasks. But this is not a robustness advantage—it’s a generalist vs specialist trade-off. For known applications, individual datasets are superior.

5.2.3 The Optimal Training Regime: Why Medium Datasets Dominate

The medium dataset superiority reflects a three-tier training regime determined by epoch count and format consistency. **Small datasets** (≤ 1 M tokens: Financial QA 0.7M, Twitter 0.28M) overtrain severely (143–352 epochs), causing memorization rather than generalization—models learn dataset-specific artifacts rather than transferable patterns. **Medium datasets** (3.6–8.5M tokens: FiQA 3.6M, FinGPT 4.1M, Alpaca 8.5M, SEC 8.1M) achieve optimal epoch counts (12–28 epochs), providing sufficient repetition for deep learning without memorization, combined with format consistency (instruction-following OR document completion, not both) that enables focused optimization. **Large datasets** (≥ 100 M tokens: News 194M, WikiText 124M) undertrain (≤ 1 epoch), providing insufficient exposure despite large vocabulary coverage, and often accumulate heterogeneous noise within single sources (News spans market analysis, earnings reports, opinion pieces with inconsistent discourse patterns).

The complete empirical hierarchy confirms these mechanisms: medium individual datasets (6.80–8.73 ppl) achieve best performance through optimal epochs \times format consistency; large mixtures (Mixed Financial 21.55 ppl) suffer from undertraining (≤ 1 epoch per component dataset) but diversity mitigates single-source noise; large individual datasets (News 32.82 ppl) combine undertraining with heterogeneous noise. Critically, Mixed Financial (21.55 ppl) OUTPERFORMS large individual News (32.82 ppl) despite both being undertrained—demonstrating that diversity helps when undertrained. However, medium individual datasets still dominate both, proving that optimal epochs + format consistency beats diversity when training dynamics are suboptimal.

The format inconsistency penalty reflects small model capacity constraints. Our 0.6B–4B parameter models lack sufficient capacity to simultaneously learn diverse formats (long-form documents + conversational Q&A + short-form tweets), vocabularies (regulatory terminology + sentiment expressions + technical jargon), and discourse patterns (narrative structure + question-answer pairs + abbreviated syntax). Format conflicts create optimization interference—loss gradients from different formats pull parameters in conflicting directions, preventing convergence to optimal solutions for any single format. Larger models (70B+ parameters) might handle mixture diversity better through increased capacity, but this defeats our edge deployment goal of lightweight, privacy-preserving models. The optimal training regime for small models prioritizes **epoch count \times format consistency \times focused diversity** over broad format coverage.

5.2.4 Domain Interference Patterns

While in-domain diversity helps, cross-domain mixing (Mixed Wiki+Financial) shows interference: Mixed Wiki+Financial achieves 26.69 ppl (4B), 24% worse than pure Mixed Financial (21.55 ppl), despite including WikiText. On WikiText specifically, the mixed approach improves performance mod-

estly compared to pure Financial, but mean financial performance degrades notably. The trade-off is unfavorable for finance-focused applications: sacrificing financial performance for a small general-domain gain.

Financial and general domains create conflicting gradients. Financial texts reward predicting domain terminology (“EBITDA” following “reported”); general texts reward different continuations (“findings” following “reported”). The model’s parameters cannot simultaneously optimize for both distributions without compromise. Mixed Wiki+Financial models average these signals, achieving moderate performance on both rather than excellence on either. The 62% variance (vs 55% pure financial) reflects this optimization conflict.

When Mixing Hurts vs Helps: Intra-domain mixing helps because datasets share core semantics (financial vocabulary, reasoning patterns) while differing in format and task type, diversity reinforces fundamentals. Cross-domain mixing hurts when domains diverge in vocabulary and reasoning (encyclopedic vs analytical), creating zero-sum trade-offs. The 50cap strategy mitigates but doesn’t eliminate interference: capping WikiText at 50% limits damage but still dilutes financial specialization. This distinction is evident comparing Table 4.2 (pure financial mixture) and Table 4.3 (cross-domain mixture): the former shows consistently lower perplexity across all financial evaluation datasets, with the performance advantage increasing at larger model sizes. Figures 4.1 and 4.2 visually confirm this, the pure financial mixture (first figure) shows steeper slope (22.6% total improvement) compared to Wiki+Financial (second figure, 15.1% improvement), indicating that domain conflict reduces scaling efficiency.

Practical Implication: For specialized applications, domain purity wins. Only mix cross-domain when explicit general-domain retention is required (e.g., conversational agents handling both financial and general queries). For finance-focused deployments, pure in-domain mixtures maximize performance.

5.2.5 Scale-Dependent Training Notes

Our experience suggests that larger models can be more sensitive to optimization settings on some datasets. While we kept LR=2e-5 for main runs, reducing LR in a handful of follow-ups helped stabilize training. We do not claim a general rule beyond this observation.

5.3 Practical Guidelines for Financial LM Pretraining

Synthesizing experimental findings into actionable recommendations:

5.3.1 Data Mixture Strategies by Use Case

Task-Specific Financial Applications (RECOMMENDED): Use individual medium datasets for optimal performance and consistency. FiQA (6.80 ppl, 19% spread) for Q&A, FinGPT (7.03 ppl, 37% spread) for sentiment, Alpaca (8.73 ppl, 11.5% spread) for instruction-following. These achieve 2.5–3.2× better perplexity AND 1.5–4.8× better consistency than mixtures. Cross-dataset tables show individual datasets excel: Alpaca achieves 6/8 boldface, FiQA 5/8, FinGPT 4/8, versus Mixed Financial 2/8. For any known application, individual datasets are superior.

Unknown Task Coverage (ONLY IF NECESSARY): Use Mixed Financial (21.55 ppl, 55% spread) ONLY when future task requirements are completely unpredictable and task coverage matters

more than optimization quality. The mixture provides baseline capability across diverse tasks but is inferior to individual datasets on both performance and consistency. This is a generalist vs specialist trade-off, not a robustness advantage. Figure 4.4 shows mixture underperformance—individual dataset curves lie substantially below.

Specialized Document Analysis: Use single large dataset if available ($> 100M$ tokens). SEC @ 4B (15.91 ppl on SEC; $\tilde{1}9\%$ relative spread across evaluations) excels for regulatory filing analysis; News @ 4B (17.47 ppl on News; $\tilde{6}6\%$ relative spread) excels for journalism. Specialization improves in-domain performance but sacrifices cross-format transfer. Figures 4.5 and 4.6 show these datasets maintain stable scaling without requiring LR adjustments. However, Tables 4.12 and 4.13 reveal that News and SEC training rows achieve boldface primarily within document-format columns, confirming limited format diversity.

For instruction-following and Q&A applications, use FiQA (3.6M tokens, 16.35 ppl) or FinGPT (4.1M tokens, 19.83 ppl) for specialized Q&A, or include in mixture for general applications. Instruction formats transfer moderately within task type ($r = 0.68 - 0.73$) but poorly to documents. The instruction-following tables (Tables 4.14 to 4.16) show boldface clustering along the diagonal and adjacent instruction rows, visualizing the format-based transfer limitation.

Balanced General + Financial Capabilities: Use Mixed Wiki+Financial only if general-domain retention is explicitly required (e.g., chatbots handling both financial and general queries). Accepts 24% financial performance cost for 16% general improvement, unfavorable for finance-focused deployments. Figure 4.2 shows reduced slope compared to pure financial mixture, and Table 4.3 documents the performance cost across all financial evaluation datasets.

Avoid: Pure WikiText for financial applications ($2.0\times$ performance degradation vs Mixed Financial on average across financial tasks), small individual datasets $< 20M$ tokens (non-viable standalone due to severe overtraining and high variability), single-format training when diverse tasks are expected (format mismatch prevents transfer). Figures 4.3, 4.10 and 4.11 provide visual evidence: WikiText requires LR adjustment and still shows poor financial transfer, while small datasets exhibit brittleness visible in both scaling curves and cross-dataset table patterns.

5.3.2 Model Size Selection

0.6B Models: Fast training (~ 6 hours for 100M tokens on Lambda Labs GPUs), low memory (4GB), suitable for rapid prototyping. Performance is acceptable for exploratory work, but variability is high (Mixed Financial: $\tilde{9}8\%$ relative spread). Use for development, experimentation, or extremely resource-constrained deployment (mobile devices).

1.7B Models: Best performance-efficiency balance. Training moderate (~ 12 hours), memory reasonable (10GB), performance strong with improved consistency vs 0.6B (Mixed Financial: $\tilde{6}3\%$ relative spread). Recommended for most applications, strong performance at substantially lower resource cost than 4B. Optimal for production deployment balancing quality and resource constraints.

4B Models: Best absolute performance (21.55 ppl, 55% relative spread) but requires careful hyperparameter tuning (LR 5×10^{-6} in affected cases) and substantial resources (20GB memory, ~ 24 hours training). Use when maximizing performance justifies cost, and when expertise for hyperparameter tuning is available. Critical: failure to tune learning rate can cause reverse scaling, one may need to reduce LR substantially at larger scales.

Scaling Decision Tree: In short, choose 0.6B under tight resource constraints (mobile, edge) and accept about a 22% loss versus 4B; pick 1.7B for the best performance-efficiency balance (about 92%

of 4B at roughly half the resources); go to 4B only when maximum performance justifies the tuning effort.

5.3.3 Learning Rate Notes

We used $\text{LR}=2\text{e-}5$ for all main runs. In three follow-ups with anomalies (WikiText, Financial QA, Twitter), smaller LRs (e.g., 1×10^{-5} or 5×10^{-6}) stabilized training. These are practical notes from our setup, not prescriptive guidelines for other contexts.

5.3.4 Token Budget Allocation

A 100M token budget proved sufficient when properly mixed, showing diminishing returns beyond that threshold for our 0.6B-4B models. Larger models ($> 7\text{B}$) may benefit from extended training (200-500M tokens), but we did not test this.

We used 50cap to prevent dominance: if the largest dataset exceeds 50% of the total, cap it there and sample others proportionally. This ensures diversity while respecting relative dataset informativeness. During training, sample each batch from the mixture distribution at the token level, not by example or epoch. Sequential exposure can cause catastrophic forgetting; token-level interleaving avoids that.

When curating datasets, format diversity comes first, followed by size (aim for $\geq 100\text{M}$ total), then quality. In-domain noisy data beats out-of-domain clean data in our runs. Don’t exclude small datasets ($< 20\text{M}$ tokens) from mixtures; they contribute valuable diversity despite non-viability as standalone sources.

5.4 Limitations and Threats to Validity

All experiments used Qwen3 (0.6B/1.7B/4B). Observations about LR behavior may be architecture- and dataset-specific. Other decoder-only transformers (LLaMA, Gemma, Phi) could behave differently; validation required. Encoder-only (BERT) or encoder-decoder (T5) models may show different mixture effects due to bidirectional attention or different pretraining objectives.

We used 50cap exclusively. Other algorithms (temperature sampling, equal mixing, DoReMi dynamic weighting) remain unexplored. The 50cap heuristic worked well but may not be optimal. Ablation studies varying cap thresholds (30%, 40%, 60%) could reveal improvements. Dynamic mixture strategies that adjust dataset weights during training based on validation loss may outperform static 50cap.

We evaluated using perplexity on held-out test sets from the same distributions as training data. This measures pretraining quality but doesn’t directly assess downstream task performance. Fine-tuned performance on financial NLP tasks (sentiment classification accuracy, Q&A F1, summarization ROUGE) may differ from pretraining perplexity rankings. Future work should validate that Mixed Financial’s pretraining advantage transfers to downstream applications.

Experiments were limited to 0.6B-4B models due to available hardware (RTX A6000 48GB, A100 40GB, H100 80GB rented from Lambda Labs). Larger models (7B, 13B, 70B) may show different patterns; mixture benefits may increase or decrease with scale. We did not investigate LR behavior beyond the few follow-ups reported here.

We systematically explored learning rates but kept other hyperparameters fixed (effective batch size

8, warmup 1000 steps, cosine schedule). Larger hyperparameter sweeps over batch size (4, 8, 16, 32), warmup ratios (1%, 3%, 5%), and schedules (linear, cosine, polynomial) may reveal better configurations. Computational budget constraints prevented exhaustive search.

Results may not generalize to other specialized domains with different characteristics. Legal text (extremely long documents, formal citations) or medical text (heavy abbreviations, multimodal integration) may show different mixture effects. The core principles (in-domain diversity, and our LR heuristics) may generalize, but specific mixture ratios and optimal configurations require domain-specific validation.

Despite these limitations, our findings provide solid empirical evidence for data mixture effects, training dynamics, and practical practices applicable to financial LM pretraining and likely informative for other specialized domains.

Chapter 6

Conclusion

This thesis investigated efficient pretraining strategies for financial language models, addressing the challenge of developing lightweight, privacy-preserving models suitable for on-device deployment. Through systematic experiments with 10 pretraining configurations across three scales (0.6B/1.7B/4B parameters), we derived empirical guidelines for data mixture composition and resource allocation, demonstrating that effective specialized models can be developed without massive computational resources while maintaining practical deployment feasibility.

6.1 Summary of Contributions

This work contributes five empirical findings and one practical deliverable, providing both empirical evidence and practical guidelines for financial language model pretraining:

6.1.1 Individual Datasets Outperform Mixtures

Contrary to conventional wisdom favoring data diversity, we found that **medium individual datasets (3.6–8.5M tokens) consistently outperform mixtures on both performance and consistency**. FiQA (6.80 ppl, 19% spread), FinGPT (7.03 ppl, 37% spread), and Alpaca (8.73 ppl, 11.5% spread) achieve 2.5–3.2× better perplexity AND 1.5–4.8× better cross-dataset consistency than Mixed Financial (21.55 ppl, 55% spread). The mixture hypothesis—that diversity improves robustness—fails empirically at fixed token budgets (100M).

This finding differs from the original hypothesis of this thesis. We designed the 50cap mixture strategy to balance large datasets (News 194M, SEC 8.1M) with smaller sources (Twitter 0.28M, FiQA 3.6M), expecting diversity to improve generalization. Instead, format inconsistency, vocabulary dilution, and multi-task interference degraded both performance and consistency. Medium individual datasets achieve superior results through focused optimization: format consistency, concentrated vocabulary exposure, and single-task training signals.

WikiText (0.6B: 9.68 ppl) shows competitive general-domain performance, but exhibits reverse scaling at 4B (31.54 ppl) due to training instability, not domain transfer failure. Including WikiText in mixtures (Mixed Wiki+Financial: 26.69 ppl, 62% spread) degrades financial performance by 24% with minimal general-domain gains—an unfavorable trade-off. Figure 4.4 shows mixture under-performance; cross-dataset tables demonstrate individual datasets capture more boldface positions (Alpaca: 6/8, FiQA: 5/8, FinGPT: 4/8) than Mixed Financial (2/8).

6.1.2 Learning Rate Notes

All primary experiments used $\text{LR}=2\text{e-}5$. In three follow-up runs that showed abnormalities (WikiText, Financial QA, Twitter), we reduced LR (e.g., to 1×10^{-5} or 5×10^{-6}) as a practical fix to stabilize training. We do not claim a general scaling rule from these observations.

6.1.3 Dataset Size Effects and Generalization

We established quantitative thresholds for dataset viability: datasets exceeding 100M tokens undertrain due to limited token budgets (≥ 1 epoch), while datasets below 1M tokens overtrain severely (143–352 epochs). Large datasets (News 194M, WikiText 124M tokens) exhibited undertraining; medium datasets (SEC 8.1M, FinGPT 4.1M, Alpaca 8.5M, FiQA 3.6M tokens) achieved optimal training dynamics (12–28 epochs) and best performance; small datasets (Financial QA 0.7M, Twitter 0.28M tokens) exhibited severe overtraining and high variability, performing well on in-distribution data but failing on out-of-distribution formats. This non-monotonic relationship reflects epoch-format consistency interactions. Medium datasets (3.6–8.5M, 12–28 epochs) achieve optimal training dynamics: sufficient repetition for deep learning without memorization, combined with format consistency that enables focused optimization. Large datasets (≥ 1 epoch) provide insufficient exposure despite large vocabulary. Small datasets (143–352 epochs) memorize rather than generalize. Large mixtures add format inconsistency (documents + Q&A + tweets)—small models (0.6B–4B) cannot simultaneously optimize for diverse formats, creating gradient conflicts that degrade performance on all tasks.

The correlation between dataset size (log transformed token count) and generalization (inverse variability) was strong ($r = -0.78$), supporting intuitions about data scale while giving concrete thresholds. Critically, small datasets remain valuable in mixtures, their contribution to format diversity and vocabulary coverage improves overall mixture quality despite standalone non-viability. The figures illustrate this: Figures 4.5 and 4.6 (large datasets) show smooth curves, while Figures 4.10 and 4.11 (small datasets) require LR interventions and exhibit erratic patterns. Cross dataset tables (Tables 4.17 and 4.18) reveal brittleness: these training rows achieve boldface only in their own columns while showing 30 to 50 ppl elsewhere.

6.1.4 Domain Transfer and Format Effects

Contrary to common assumptions that domain vocabulary drives transfer, we found format consistency determines generalization more than semantic domain. Long-form financial documents (News \leftrightarrow SEC) exhibited strongest transfer ($r = 0.82$ cross-perplexity correlation), while instruction-format transfers (FiQA \leftrightarrow FinGPT \leftrightarrow Alpaca) achieved moderate correlation ($r = 0.68$ – 0.73). Cross-format transfer failed: document-pretrained models achieved 2–3 \times worse perplexity on instruction tasks (and vice versa) despite shared financial vocabulary.

Domain transfer proved asymmetric: financial pretraining enabled reasonable general-domain performance (WikiText perplexity competitive with specialized general-domain models), but general pretraining failed catastrophically for financial tasks (WikiText pretraining: 41.96 mean financial ppl vs 21.55 for Mixed Financial). This asymmetry reflects vocabulary coverage, financial text includes substantial general vocabulary, but general corpora lack domain-specific terminology (EBITDA, prospectus, liquidity ratios).

These findings suggest **format diversity should be prioritized over domain purity** when

curating pretraining mixtures. A mixture spanning documents, dialogues, and Q&A formats within the financial domain generalizes better than narrow focus on a single format, even with larger data volume. Cross-dataset comparison tables provide striking visual evidence: Tables 4.12 and 4.13 show boldface clustering along the News-SEC diagonal (long-form transfer), Tables 4.14 to 4.16 exhibit diagonal boldface patterns plus adjacency (instruction-format clustering), while Table 4.17 shows complete isolation (boldface only in Twitter’s own column). These format-driven transfer patterns explain why medium individual datasets outperform diverse mixtures: format consistency enables focused optimization within small model capacity constraints, while format diversity creates optimization conflicts that small models cannot resolve.

6.1.5 Model Size Selection for Resource-Constrained Settings

We demonstrated that **1.7B models offer optimal performance-efficiency balance**, achieving 92% of 4B’s performance (24.12 vs 21.55 ppl on Mixed Financial) while requiring 50% memory (10GB vs 20GB) and 50% training time (12 vs 24 hours for 100M tokens on Lambda Labs GPUs). This finding directly addresses the thesis motivation of developing lightweight, privacy-preserving models for on-device deployment.

0.6B models remain viable for rapid prototyping and extremely resource-constrained scenarios (mobile devices), accepting 22% performance degradation (27.84 ppl) for 3 \times faster training. 4B models justify their cost only when maximizing absolute performance is critical and hyperparameter tuning expertise is available, improper learning rate selection at this scale causes training collapse, making 4B models less stable than smaller alternatives.

For the privacy-preserving financial chatbot application motivating this thesis, **1.7B represents the recommended deployment target**: small enough for laptop inference on consumer hardware, large enough for acceptable performance, and stable enough for reliable training without extensive hyperparameter search.

6.1.6 Open-Source Reproducible Pipeline

Beyond empirical findings, we delivered a production-ready training pipeline supporting 26 datasets (10 financial classification, 11 generative Q&A, 5 pretraining corpora), multiple model architectures (Qwen, LLaMA, Gemma, Phi), and advanced techniques (LoRA, FlashAttention, MOE support). The pipeline includes automatic experiment naming, TensorBoard logging, checkpoint management, and comprehensive documentation, lowering barriers to entry for researchers.

All experiments in this thesis are fully reproducible using documented commands and publicly available datasets. The codebase has been structured for extensibility, adding new datasets requires minimal modifications (label mappings, prompt templates). This contribution addresses reproducibility challenges in NLP research and provides a foundation for future work in specialized domain adaptation.

6.2 Implications for Practice and Research

6.2.1 Actionable Deployment Guidelines

Financial institutions and fintech companies developing on-premise NLP systems can directly apply our findings:

Data Strategy: Use individual medium datasets (3.6–8.5M tokens) for optimal performance and consistency. FiQA for Q&A (6.80 ppl, 19% spread), FinGPT for sentiment (7.03 ppl, 37% spread), Alpaca for instruction-following (8.73 ppl, 11.5% spread). These achieve 2.5–3.2 \times better perplexity and 1.5–4.8 \times better consistency than mixtures. Only use mixtures when future task requirements are completely unknown and task coverage matters more than optimization quality. Avoid general-domain corpora (WikiText, C4) unless explicitly required—they show competitive performance at small scales (0.6B) but unstable scaling.

Model Selection: Deploy 1.7B models for production applications balancing quality and resources. Use 0.6B for prototyping and testing. Reserve 4B for scenarios where performance justifies cost and expertise for careful LR tuning is available.

Training Configuration: Use LR=2e-5 for main runs; if a larger model shows instability on a dataset, try a smaller LR (e.g., 1×10^{-5} or 5×10^{-6}). Monitor loss curves to detect instability. Allocate 100M token budgets, diminishing returns beyond this threshold for sub-5B models. Use standard configurations (AdamW, cosine schedules, 1000 warmup steps) which proved stable across experiments.

Privacy and Compliance: On-device deployment with 1.7B models enables GDPR-compliant financial NLP without data transmission to external services, critical for banks, investment firms, and financial advisors handling sensitive client information.

6.2.2 For Researchers: Open Questions and Methodological Lessons

Our work highlights underexplored areas in LM scaling research:

Hyperparameter Sensitivity: Scaling laws literature (J. Kaplan et al. 2020; Hoffmann et al. 2022) focuses on compute-optimal training but provides limited guidance for hyperparameter transfer across scales. We did not develop or test a learning-rate scaling theory. Future work should investigate when simple LR reductions help, and explore principled relationships for batch size, warmup steps, weight decay, and optimizer hyperparameters.

Domain Adaptation Theory: Why does format dominate vocabulary for transfer? Our findings challenge intuitions about domain similarity. Theoretical frameworks explaining when syntactic structure (format) versus semantic content (vocabulary) drives generalization would inform curriculum design and transfer learning strategies. Neuroscience-inspired probing of intermediate representations may reveal whether format information resides in different layers/attention heads than vocabulary.

Data Mixing Algorithms: We used static 50cap throughout. Dynamic strategies (DoReMi, temperature sampling, curriculum learning) that adjust mixture weights based on validation loss or task difficulty may outperform static mixing. Ablation studies varying cap thresholds (30%, 40%, 60%) would clarify sensitivity. Meta-learning approaches that optimize mixture ratios as hyperparameters represent promising future directions.

Evaluation Methodology: We assessed pretraining quality via perplexity on held-out test sets.

Downstream task evaluation (sentiment classification accuracy, Q&A F1, summarization ROUGE) would validate that pretraining improvements transfer to practical applications. Establishing correlations between pretraining perplexity and downstream performance across diverse tasks would enable efficient model selection without exhaustive downstream evaluation.

6.2.3 For Industry: Privacy-Preserving Financial AI

This work directly addresses emerging regulatory and business needs:

Regulatory Compliance: GDPR, CCPA, and emerging AI regulations increasingly prohibit or restrict transmission of financial data to external APIs (OpenAI, Anthropic). On-device models trained using our methods enable compliant NLP for document analysis, risk assessment, and customer service without data exfiltration.

Competitive Differentiation: Financial institutions accumulating proprietary datasets (transaction records, analyst reports, client communications) can leverage specialized pretraining to develop competitive advantages, custom models trained on confidential data without exposing information to vendors.

Cost Efficiency: Cloud API costs for LLM inference (\$0.002 – 0.03 per 1K tokens for GPT-4 class models) accumulate rapidly at scale. On-premise 1.7B models reduce marginal costs to negligible levels (electricity, amortized hardware), enabling aggressive deployment for high-volume applications (transaction categorization, automated report generation).

Latency and Reliability: Local inference eliminates network latency and dependency on external service availability, critical for real-time trading applications and customer-facing systems requiring $\leq 100\text{ms}$ response times.

6.3 Future Research Directions

6.3.1 Scaling to Larger Models and Architectures

Our experiments covered 0.6B-4B parameters ($6.7\times$ range) on Qwen3 architecture exclusively. Critical open questions:

Larger Scales: Do 7B, 13B, 70B models exhibit the same mixture effects and similar sensitivity to optimization settings? Larger models may benefit more from diverse pretraining (improved few-shot generalization) or less (stronger preexisting representations). Understanding learning-rate behavior at these scales requires focused empirical and theoretical work.

Architectural Diversity: LLaMA, Gemma, Mistral, Phi use different architectural choices (grouped-query attention variants, different activation functions, rotary embeddings). Validating our findings across architectures would establish generality. Encoder-decoder models (T5, BART) and encoder-only models (BERT, RoBERTa) may show different mixture effects due to bidirectional attention or different pretraining objectives (masked language modeling vs causal).

MOE Architectures: Mixture-of-Experts models (Mixtral, DeepSeek-MOE) offer computational efficiency through sparse activation. Do MOE models benefit more from data mixtures (experts specializing on subdomains) or less (already mixture-like internally)? MOE-specific mixture strategies matching expert count to dataset count represent unexplored territory.

6.3.2 Advanced Mixture Optimization

Our 50cap strategy worked well but wasn't rigorously optimized. Future work should explore:

Dynamic Mixture Schedules: Curriculum learning approaches that shift mixture composition during training, start with general data for basic capabilities, transition to specialized data for domain expertise. Or reverse: start specialized to establish domain vocabulary, add general data to improve stability.

Adaptive Weighting: Use validation loss gradients or uncertainty estimates to identify which datasets currently contribute most to learning. Upweight informative datasets, downweight exhausted sources. DoReMi (Xie et al. 2023) provides reference implementation; adaptation to specialized domains requires experimentation.

Mixture Ablations: Systematically vary cap threshold (30%, 40%, 50%, 60%, 70%) and measure sensitivity. Optimal threshold may depend on dataset size distribution, highly imbalanced mixtures (one dataset 90% of total) may require aggressive capping, while balanced mixtures may prefer minimal intervention.

Multi-Stage Mixing: Train separate models on individual datasets, then merge via model averaging, task arithmetic, or TIES merging. Compare to simultaneous mixture training. Sequential training (pretrain on Dataset A, continue on Dataset B) versus concurrent mixing represents under-explored design space.

6.3.3 Comprehensive Downstream Evaluation

This thesis assessed pretraining quality via perplexity. Validation on downstream applications would strengthen practical relevance:

Financial Sentiment Analysis: FPB, FiQA-SA, Twitter Financial sentiment datasets. Compare finetuning performance from different pretrained checkpoints (Mixed Financial vs WikiText vs single-dataset). Measure few-shot and zero-shot transfer.

Financial Q&A: FinQA, ConvFinQA, Alpaca-Finance benchmarks. Assess both extractive (span selection) and generative (free-form answer) settings. Evaluate factual accuracy and hallucination rates, critical for financial applications.

Document Summarization: SEC filing summarization, earnings call summarization. Metrics: ROUGE, BERTScore, human evaluation for factuality and conciseness. Privacy-preserving summarization represents high-value application for on-device models.

Long-Context Understanding: Financial documents often exceed 10K tokens (10-K filings, prospectuses). Evaluate long-context capabilities using retrieval-augmented generation or extended-context versions of Qwen3 (32K+ tokens). Does mixture pretraining improve long-document coherence?

6.3.4 Multi-Stage Pretraining Strategies

Our single-stage approach (pretrain directly on financial mixtures) represents one point in a broader design space:

General → Domain Adaptation: Pretrain on WikiText or C4 for general capabilities, then continue pretraining on financial data. Compare to direct financial pretraining. Theory suggests general stage builds stable syntax/reasoning, domain stage adds specialized vocabulary, empirical

validation needed.

Domain → Task Specialization: Pretrain on broad financial mixture, then continue on task-specific data (e.g., only sentiment data for sentiment model). Balances general financial knowledge with task-specific optimization.

Mixture schedules. Gradually shift mixture composition across training, start balanced, progressively upweight high-priority datasets. Or inverse: start specialized, progressively diversify to improve stability.

Optimal strategies likely depend on target application and available data. Decision frameworks are needed: "If you have 10M domain tokens and 100M general tokens, use Strategy X; if 100M domain and 10M general, use Strategy Y."

6.3.5 Open Questions

We did not study learning-rate theory or its interaction with model size and batch size. Understanding when simple LR reductions help, and when they do not, remains open work, alongside broader questions on mixture schedules and transfer to downstream tasks.

6.4 Closing Remarks

This thesis shows that effective specialized language models can be developed without massive computational resources or diverse data mixtures. By selecting focused medium datasets (3.6–8.5M tokens), using stable training settings, and targeting lightweight 1–2B parameter models, one can train privacy-preserving financial NLP systems suitable for on-device deployment. Contrary to expectations, individual datasets (FiQA, FinGPT, Alpaca) consistently outperform mixtures on both performance (2.5–3.2× better) and consistency (1.5–4.8× better). Simple focused choices outperform complex diverse strategies.

The core insight—that **individual medium datasets (3.6–8.5M tokens) consistently outperform mixtures on both performance and consistency**—challenges conventional wisdom favoring data diversity. At fixed token budgets (100M), FiQA/FinGPT/Alpaca achieve 2.5–3.2× better perplexity and 1.5–4.8× better consistency than 7-dataset mixtures. Format inconsistency, vocabulary dilution, and multi-task interference degrade mixture performance despite anticipated diversity benefits. This finding has important implications: specialized pretraining should prioritize focused, high-quality medium datasets over diverse mixtures. For domains with curated data (finance, legal, medical), individual dataset optimization offers superior performance at lower cost than either mixtures or general-purpose model adaptation.

As privacy regulations tighten and organizations recognize competitive value in proprietary data, on-device specialized models will become increasingly important. This work provides empirical foundations and practical guidelines for developing such systems, contributing to a future where powerful NLP capabilities are accessible without sacrificing privacy, incurring ongoing API costs, or depending on external service providers.

The open-source pipeline and reproducible experimental methodology lower barriers to entry, enabling researchers to build on these findings and extend them to new domains, architectures, and applications. By sharing not just results but complete implementations, we hope to accelerate progress toward privacy-preserving, efficient, and democratically accessible language understanding

for specialized domains.

Bibliography

- Aharoni, Roee and Yoav Goldberg (2020). “Unsupervised Domain Clusters in Pretrained Language Models”. In: *arXiv preprint arXiv:2004.02105*. URL: <https://arxiv.org/abs/2004.02105>.
- Araci, Dogu (2019). “Finbert: Financial sentiment analysis with pre-trained language models”. In: *arXiv preprint arXiv:1908.10063*.
- Arivazhagan, Naveen, Ankur Bapna, Orhan Firat, Dmitry Lepikhin, Melvin Johnson, Maxim Krikun, Mia Xu Chen, Yuan Cao, George F. Foster, Colin Cherry, Wolfgang Macherey, Zhifeng Chen, and Yonghui Wu (2019). “Massively Multilingual Neural Machine Translation in the Wild: Findings and Challenges”. In: *arXiv preprint arXiv:1907.05019*. URL: <http://arxiv.org/abs/1907.05019>.
- Bai, Jinze, Shuai Bai, Yunfei Chu, Zeyu Cui, Kai Dang, Xiaodong Deng, Yang Fan, Wenbin Ge, Yu Han, Fei Huang, Binyuan Hui, Luo Ji, Mei Li, Junyang Lin, Runji Lin, Dayiheng Liu, Gao Liu, Chengqiang Lu, Keming Lu, Jianxin Ma, Rui Men, Xingzhang Ren, Xuancheng Ren, Chuanqi Tan, Sinan Tan, Jianhong Tu, Peng Wang, Shijie Wang, Wei Wang, Shengguang Wu, Benfeng Xu, Jin Xu, An Yang, Hao Yang, Jian Yang, Shusheng Yang, Yang Yao, Bowen Yu, Hongyi Yuan, Zheng Yuan, Jianwei Zhang, Xingxuan Zhang, Yichang Zhang, Zhenru Zhang, Chang Zhou, Jingren Zhou, Xiaohuan Zhou, and Tianhang Zhu (2023). “Qwen Technical Report”. In: *arXiv preprint arXiv:2309.16609*.
- Bengio, Yoshua, Jérôme Louradour, Ronan Collobert, and Jason Weston (2009). “Curriculum learning”. In: *Proceedings of the 26th Annual International Conference on Machine Learning*. ACM, pp. 41–48. DOI: [10.1145/1553374.1553380](https://doi.org/10.1145/1553374.1553380).
- Brown, Tom, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, et al. (2020). “Language models are few-shot learners”. In: *Advances in neural information processing systems* 33, pp. 1877–1901.
- Chen, Zhiyu, Wenhui Chen, Charese Smile, Sameena Shah, Iana Borova, Dylan Langdon, Reema Moussa, Matt Beane, Ting-Hao Huang, Bryan Routledge, and William Yang Wang (2021). “FinQA: A Dataset of Numerical Reasoning over Financial Data”. In: *arXiv preprint arXiv:2109.00122*. URL: <https://arxiv.org/abs/2109.00122>.
- Devlin, Jacob, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova (2019). “BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding”. In: *Proceedings of NAACL-HLT 2019*, pp. 4171–4186. DOI: [10.18653/v1/n19-1423](https://doi.org/10.18653/v1/n19-1423). URL: <https://doi.org/10.18653/v1/n19-1423>.
- French, Robert M (1999). “Catastrophic forgetting in connectionist networks”. In: *Trends in Cognitive Sciences* 3.4, pp. 128–135. DOI: [10.1016/S1364-6613\(99\)01294-2](https://doi.org/10.1016/S1364-6613(99)01294-2).

- Gao, Leo, Stella Biderman, Sidney Black, Laurence Anthony, Xenia Golding, Horace Hoppe, Connor Foster, Jason Phang, Anish He, Aman Thite, Andy Nabeshima, Shawn Presser, and Connor Leahy (2021). “The Pile: An 800GB Dataset of Diverse Text for Language Modeling”. In: *arXiv preprint arXiv:2101.00027*. URL: <https://arxiv.org/abs/2101.00027>.
- Gururangan, Suchin, Ana Marasović, Swabha Swamyamdipta, Kyle Lo, Iz Beltagy, Doug Downey, and Noah A Smith (2020). “Don’t stop pretraining: Adapt language models to domains and tasks”. In: *arXiv preprint arXiv:2004.10964*.
- Hoffmann, Jordan, Sebastian Borgeaud, Arthur Mensch, Elena Buchatskaya, Trevor Cai, Eliza Rutherford, Diego de Las Casas, Lisa Anne Hendricks, Johannes Welbl, Aidan Clark, et al. (2022). “Training compute-optimal large language models”. In: *arXiv preprint arXiv:2203.15556*.
- Hu, Edward J., Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen (2021). “LoRA: Low-Rank Adaptation of Large Language Models”. In: *arXiv preprint arXiv:2106.09685*. URL: <https://arxiv.org/abs/2106.09685>.
- Huang, Allen H., Hui Wang, and Yi Yang (2023). “FinBERT: A Large Language Model for Extracting Information from Financial Text”. In: *Contemporary Accounting Research* 40.2, pp. 806–841. DOI: 10.1111/1911-3846.12832.
- Javaheripi, Mojtaba, Sébastien Bubeck, Marah Abdin, Jyoti Aneja, Sébastien Bubeck, Caio César Teodoro Mendes, Weizhu Chen, Allie Del Giorno, Ronen Eldan, Sivakanth Gopi, et al. (2023). *Phi-2: The surprising power of small language models*. Microsoft Research Blog. URL: <https://www.microsoft.com/en-us/research/blog/phi-2-the-surprising-power-of-small-language-models/>.
- Kaplan, Jared, Sam McCandlish, Tom Henighan, Tom B Brown, Benjamin Chess, Rewon Child, Scott Gray, Alec Radford, Jeffrey Wu, and Dario Amodei (2020). “Scaling laws for neural language models”. In: *arXiv preprint arXiv:2001.08361*.
- Kingma, Diederik P. and Jimmy Ba (2014). “Adam: A Method for Stochastic Optimization”. In: *arXiv preprint arXiv:1412.6980*. URL: <https://arxiv.org/abs/1412.6980>.
- Kirkpatrick, James, Razvan Pascanu, Neil Rabinowitz, Joel Veness, Guillaume Desjardins, Andrei A Rusu, Kieran Milan, John Quan, Tiago Ramalho, Agnieszka Grabska-Barwinska, Demis Hassabis, Claudia Clopath, Dharshan Kumaran, and Raia Hadsell (2017). “Overcoming catastrophic forgetting in neural networks”. In: *Proceedings of the National Academy of Sciences* 114.13, pp. 3521–3526. DOI: 10.1073/pnas.1611835114.
- Lee, Yoonho, Annie S. Chen, Fahim Tajwar, Ananya Kumar, Huaxiu Yao, Percy Liang, and Chelsea Finn (2022). “Surgical Fine-Tuning Improves Adaptation to Distribution Shifts”. In: *arXiv preprint arXiv:2210.11466*. URL: <https://arxiv.org/abs/2210.11466>.
- Longpre, Shayne, Yao Hou, Aakanksha Deshpande, He He, Thibault Sellam, Alex Tamkin, Slav Petrov, Denny Zhou, Jason Wei, Yi Tay, Quoc V. Le, et al. (2023). “A Pretrainer’s Guide to Training Data: Measuring the Effects of Data Age, Domain Coverage, Quality, & Toxicity”. In: *arXiv preprint arXiv:2305.13169*. URL: <https://arxiv.org/abs/2305.13169>.
- McCloskey, Michael and Neal J. Cohen (1989). “Catastrophic Interference in Connectionist Networks: The Sequential Learning Problem”. In: *Psychology of Learning and Motivation*. Elsevier, pp. 109–165. DOI: 10.1016/S0079-7421(08)60536-8.

- Merity, Stephen, Caiming Xiong, James Bradbury, and Richard Socher (2017). “Pointer sentinel mixture models”. In: *International Conference on Learning Representations*.
- Mitra, Arindam, Luciano Del Corro, Shweta Mahajan, Andres Codas, Clarisse Simoes, Sahaj Agarwal, Xuxi Chen, Anastasia Razdaibiedina, Erik Jones, Kriti Aggarwal, Hamid Palangi, Guoqing Zheng, Corby Rosset, Hamed Khanpour, and Ahmed Awadallah (2023). “Orca 2: Teaching Small Language Models How to Reason”. In: *arXiv preprint arXiv:2311.11045*. URL: <https://arxiv.org/abs/2311.11045>.
- Narayanan, Deepak, Mohammad Shoeybi, Jared Casper, Patrick LeGresley, Mostafa Patwary, Vijay Anand Korthikanti, Dmitri Vainbrand, Preethvi Kashinkunti, Julie Bernauer, Bryan Catanzaro, Amar Phanishayee, and Matei Zaharia (2021). “Efficient Large-Scale Language Model Training on GPU Clusters Using Megatron-LM”. In: *arXiv preprint arXiv:2104.04473*. URL: <https://arxiv.org/abs/2104.04473>.
- Pan, Sinno Jialin and Qiang Yang (2010). “A Survey on Transfer Learning”. In: *IEEE Transactions on Knowledge and Data Engineering*. Vol. 22, pp. 1345–1359. DOI: 10.1109/TKDE.2009.191.
- Quiñonero-Candela, Joaquin, Masashi Sugiyama, Anton Schwaighofer, and Neil D. Lawrence, eds. (2008). *Dataset Shift in Machine Learning*. MIT Press. DOI: 10.7551/mitpress/9780262170055.001.0001. URL: <https://doi.org/10.7551/mitpress/9780262170055.001.0001>.
- Radford, Alec, Jeffrey Wu, Rewon Child, David Luan, Dario Amodei, Ilya Sutskever, et al. (2019). “Language models are unsupervised multitask learners”. In: *OpenAI blog* 1.8, p. 9.
- Raffel, Colin, Noam Shazeer, Adam Roberts, Katherine Lee, Sharan Narang, Michael Matena, Yanqi Zhou, Wei Li, and Peter J. Liu (2020). “Exploring the Limits of Transfer Learning with a Unified Text-to-Text Transformer”. In: *Journal of Machine Learning Research* 21, 140:1–140:67. URL: <https://jmlr.org/papers/v21/20-074.html>.
- Rajbhandari, Samyam, Jeff Rasley, Olatunji Ruwase, and Yuxiong He (2020). “ZeRO: Memory optimizations Toward Training Trillion Parameter Models”. In: *SC20: International Conference for High Performance Computing, Networking, Storage and Analysis*. IEEE, pp. 1–16. DOI: 10.1109/SC41405.2020.00024. URL: <https://doi.org/10.1109/SC41405.2020.00024>.
- Regulation (EU) 2016/679 of the European Parliament and of the Council of 27 April 2016 on the protection of natural persons with regard to the processing of personal data and on the free movement of such data (General Data Protection Regulation)* (2016). Official Journal of the European Union. URL: <https://eur-lex.europa.eu/eli/reg/2016/679/oj>.
- Sanh, Victor, Albert Webson, Colin Raffel, Stephen H. Bach, Lintang Sutawika, Zaid Alyafeai, Antoine Chaffin, Arnaud Stiegler, Teven Le Scao, Arun Raja, et al. (2022). “Multitask Prompted Training Enables Zero-Shot Task Generalization”. In: *arXiv preprint arXiv:2110.08207*. URL: <https://arxiv.org/abs/2110.08207>.
- Tay, Yi, Mostafa Dehghani, Vinh Q. Tran, Xavier Garcia, Jason Wei, Xuezhi Wang, Hyung Won Chung, Siamak Shakeri, Dara Bahri, Tal Schuster, Huaixiu Steven Zheng, Denny Zhou, Neil Houlsby, and Donald Metzler (2022). “UL2: Unifying Language Learning Paradigms”. In: *arXiv preprint arXiv:2205.05131*. URL: <https://arxiv.org/abs/2205.05131>.
- Team, Gemma, Thomas Mesnard, Cassidy Hardin, Robert Dadashi, Surya Bhupatiraju, Shreya Pathak, Laurent Sifre, Morgane Rivière, Mihir Sanjay Kale, Juliette Love, et al. (2024). *Gemma: Open Models Based on Gemini Research and Technology*. URL: <https://arxiv.org/abs/2403.08295>.

- Touvron, Hugo, Thibaut Lavril, Gautier Izacard, Xavier Martinet, Marie-Anne Lachaux, Timothée Lacroix, Baptiste Rozière, Naman Goyal, Eric Hambro, Faisal Azhar, et al. (2023). “Llama: Open and efficient foundation language models”. In: *arXiv preprint arXiv:2302.13971*.
- Vaswani, Ashish, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin (2017). “Attention is All you Need”. In: *Advances in Neural Information Processing Systems 30*, pp. 5998–6008. URL: <https://proceedings.neurips.cc/paper/2017/hash/3f5ee243547dee91fdb053c1c4a845aa-Abstract.html>.
- Wu, Shijie, Ozan Irsoy, Steven Lu, Vadim Dabrowski, Mark Dredze, Sebastian Gehrmann, Prabhanjan Kambadur, David S. Rosenberg, and Gideon Mann (2023). “BloombergGPT: A Large Language Model for Finance”. In: *arXiv preprint arXiv:2303.17564*. URL: <https://arxiv.org/abs/2303.17564>.
- Xia, Mengzhou, Tianyu Gao, Zhiyuan Zeng, and Danqi Chen (2023). “Sheared llama: Accelerating language model pre-training via structured pruning”. In: *arXiv preprint arXiv:2310.06694*.
- Xie, Sang Michael, Hieu Pham, Xuanyi Dong, Nan Du, Hanxiao Liu, Yifeng Lu, Percy Liang, Quoc V. Le, Tengyu Ma, and Adams Wei Yu (2023). “DoReMi: Optimizing Data Mixtures Speeds Up Language Model Pretraining”. In: *arXiv preprint arXiv:2305.10429*. URL: <https://arxiv.org/abs/2305.10429>.
- Yang, An, Baosong Yang, Binyuan Hui, Bo Zheng, Bowen Yu, Chang Zhou, Chengpeng Li, Chengyuan Li, Dayiheng Liu, Fei Huang, et al. (2024). “Qwen2 Technical Report”. In: *arXiv preprint arXiv:2407.10671*.
- Yang, Hongyang, Xiao-Yang Liu, and Christina Dan Wang (2023). “FinGPT: Open-Source Financial Large Language Models”. In: *arXiv preprint arXiv:2306.06031*. URL: <https://arxiv.org/abs/2306.06031>.
- Yang, Yi, Mark Christopher Siy UY, and Allen Huang (2020). “FinBERT: A Pretrained Language Model for Financial Communications”. In: *arXiv preprint arXiv:2006.08097*. URL: <https://arxiv.org/abs/2006.08097>.
- Zhang, Susan, Stephen Roller, Naman Goyal, Mikel Artetxe, Moya Chen, Shuhui Chen, Christopher Dewan, Mona Diab, Xian Li, Xi Victoria Lin, Todor Mihaylov, Myle Ott, Sam Shleifer, Kurt Shuster, Daniel Simig, Punit Singh Koura, Anjali Sridhar, Tianlu Wang, and Luke Zettlemoyer (2022). “OPT: Open Pre-trained Transformer Language Models”. In: *arXiv preprint arXiv:2205.01068*. URL: <https://arxiv.org/abs/2205.01068>.
- Zhuang, Fuzhen, Zhiyuan Qi, Keyu Duan, Dongbo Xi, Yongchun Zhu, Hengshu Zhu, Hui Xiong, and Qing He (2021). “A Comprehensive Survey on Transfer Learning”. In: *Proceedings of the IEEE* 109, pp. 43–76. DOI: [10.1109/JPROC.2020.3004555](https://doi.org/10.1109/JPROC.2020.3004555).

Eidesstattliche Erklärung

Der/Die Verfasser/in erklärt an Eides statt, dass er/sie die vorliegende Arbeit selbstständig, ohne fremde Hilfe und ohne Benutzung anderer als die angegebenen Hilfsmittel angefertigt hat. Die aus fremden Quellen (einschliesslich elektronischer Quellen) direkt oder indirekt übernommenen Gedanken sind ausnahmslos als solche kenntlich gemacht. Die Arbeit ist in gleicher oder ähnlicher Form oder auszugsweise im Rahmen einer anderen Prüfung noch nicht vorgelegt worden.

.....
Ort, Datum

.....
Unterschrift des/der Verfassers/in